The Window for Preventive Action Before AI Saturates Most Cognitive Benchmarks is Closing

A Preprint

Jérémy Andréoletti* **Antoine Maier** **Laura Domenech** **Tom David**

General-Purpose AI Policy Lab

Abstract

This position paper argues that without coordinated intervention, AI systems will by default saturate most existing cognitive benchmarks by 2030, approaching or exceeding human expert performance across a broad range of measurable tasks. Using a joint hierarchical Bayesian model, we forecast frontier performance trajectories across 98 benchmarks spanning reasoning, mathematics, coding, cybersecurity, scientific knowledge, and agentic capabilities. We find that nearly all current benchmarks (98% in our set, 80% CI: 96–99%) are on track to saturate before 2030. This result holds for all 8 model variants and three saturation thresholds we tested. Security-critical benchmarks (cybersecurity, autonomous AI R&D, biology, chemistry) saturate earlier still, all before 2028. Benchmark saturation is not artificial general intelligence, but these results add to the body of evidence that superhuman performance on cognitive tasks is coming sooner than commonly assumed. Acknowledging that we neither understand nor control the behavior of current general-purpose AI (GPAI) models, this timeline leaves limited runway before irreversible impacts. We outline implications for global governance and international coordination, AI safety research, and evaluation practices.

Keywords AI progress · benchmark saturation · capability forecasting · AI safety · AGI

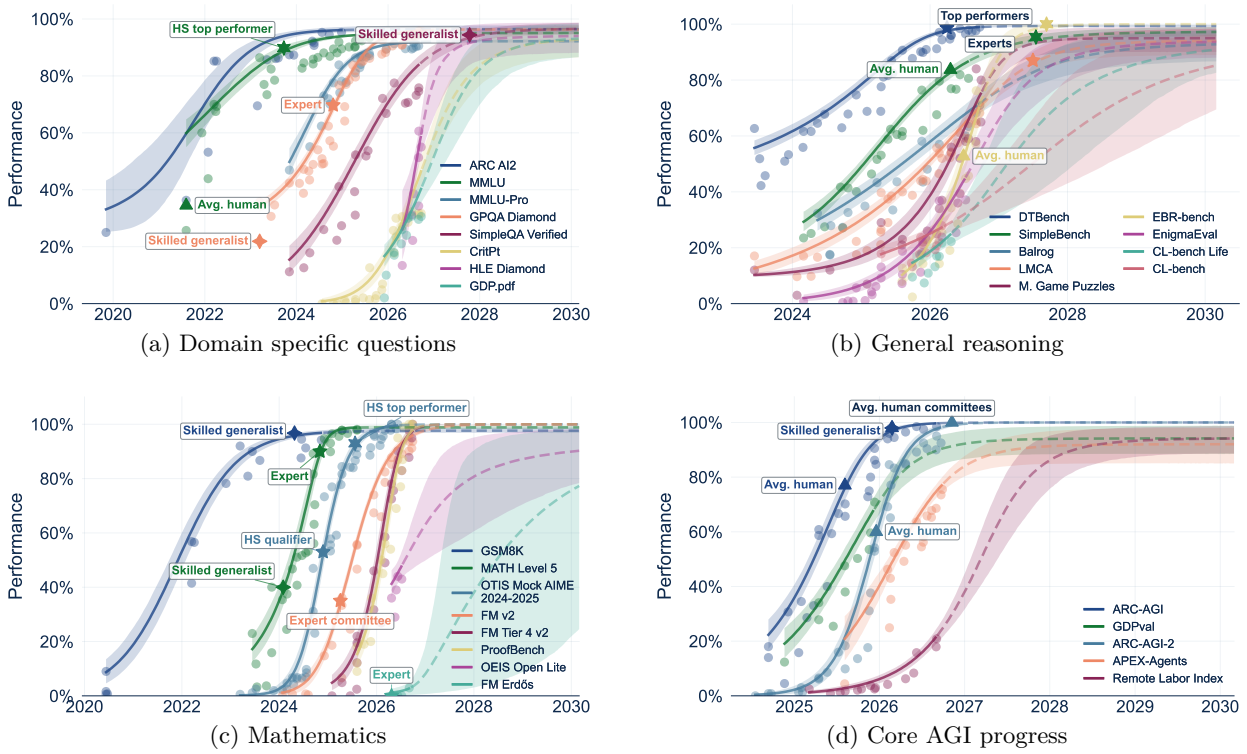
1 Introduction

The pace at which new benchmarks have been developed over the past five years has only been matched by the speed at which these benchmarks have been climbed by each new generation of General-Purpose AI (GPAI) models (Epoch AI (2024); Bengio et al. (2025); AISI (2025); Figure 1). Doctorate-level science questions, competition mathematics and autonomous software engineering tasks, long considered out of reach, are now routinely solved by frontier models. Yet expert surveys often place human-level performance on broad cognitive tasks decades away (Grace et al., 2018, 2025), or consider its timing highly uncertain.

Our position: Based on an analysis of 98 benchmark trajectories, we argue that AI systems are on track to saturate most existing cognitive benchmarks by 2030, approaching or exceeding human expert performance across a wide range of cognitive tasks. This projection, which includes security-critical capabilities, leaves a short window for preventive action before irreversible impacts. This is not a prediction that AGI will necessarily arrive by 2030, as the relationship between benchmark performance and general intelligence remains contested (Chollet, 2019). However, it is a call to take seriously the empirical trend that AI is saturating our best evaluations faster than anticipated, and faster than current coordination efforts for global GPAI risk management.

If correct, this timeline leaves little time to develop robust alignment or control techniques, given the current lack of any reliable method for understanding and steering powerful AI systems (Casper et al., 2023; Bengio et al., 2025; Maier et al., 2025; Ngo et al., 2025). It also implies an urgency regarding international coordination on AI; the earlier multilateral discussions gain momentum, the more time will be available for countries

*Corresponding author: jeremy.andreoletti@gmail.com



(a) Domain specific questions (b) General reasoning

(c) Mathematics (d) Core AGI progress

Figure 1: **Benchmark trajectories across capability categories.** Points show frontier model scores at release; solid lines show posterior mean trajectories; shaded regions indicate 80% credible intervals. Dashed lines extrapolate to 2030. Human reference points are star-like symbols whose number of branches denotes the expertise level: 3 for average human, 4 for skilled generalist, 5 for domain expert, and 6 for top performer. See more details in Appendix F. All categories show rapid progress towards saturation, with several benchmarks already exceeding human expert baselines.

to converge on a global course of action before GPAI models could start posing irreversible security issues. Even in optimistic safety scenarios, this pace of progress leaves only a few years for institutions to adapt to transformative AI capabilities.

This position therefore chains three claims resting on evidence of decreasing strength. **(C1) Saturation:** frontier models reach 95% of the score range on nearly all benchmarks in our set before 2030. This is the output of our forecasts, robust to every modeling choice we varied (Appendix B) and validated by temporal holdout (Appendix C). **(C2) Human-expert parity:** saturation implies matching or exceeding domain-expert performance on the measured tasks (42 of our 98 benchmarks carry human baselines, including domain-expert baselines for 21 of them, see Appendix F). **(C3) A short window before large impacts:** reaching that level without reliable means of controlling these systems leaves limited time to act before consequences become hard to reverse. This argument, based on the capabilities involved, is detailed in Section 4.4.

Our contribution is threefold: (1) We provide quantitative evidence from 98 benchmarks and 56 human baselines showing convergent saturation trajectories based on a new modeling framework; (2) We analyze AI progress specifically in security-critical capability domains; (3) We outline implications for researchers and policymakers.

2 Evidence: Benchmark Trajectories

2.1 Data and Methodology

We analyze performance trajectories on 98 benchmarks spanning diverse cognitive capabilities: commonsense reasoning, scientific knowledge, mathematical problem-solving, code generation, cybersecurity, agentic computer use, and multimodal understanding. Benchmark scores are sourced from the Epoch AI Benchmark database (Epoch AI, 2024), including its compilation of cybersecurity results from the UK AI Security Institute and other developers (Chauvin et al., 2026), Scale AI leaderboards (Scale AI, 2025b), evaluations by the RAND Corporation (Dev et al., 2025), Kaggle Open Benchmarks (Kaggle, 2026) and the Cybench leaderboard (Zhang et al., 2025). We exclude benchmarks whose scores cover too small a share of the frontier models of their period, deprecated series, and series with documented quality problems. All included benchmarks are described and cited in Appendix E.1. To ensure that our projections are based on a well-defined observation window and can be reproduced, we restrict the fitting data to models released up to October 1, 2026. We contextualize these projections by overlaying human performance baselines where available. These baselines range from average crowdworkers to world-class experts, providing interpretable reference points for AI progress. Detailed human baseline sources are provided in Appendix F.

Terminology. We use *score* for a published benchmark result and *capability* for the best performance a model can exhibit across elicitation conditions, given this is where tail risks should come from. *Saturation* means reaching 95% of the range between random chance and the estimated asymptote.

To project future performance, we employ hierarchical Bayesian models which capture the S-shaped trajectories empirically observed in benchmark progress. Our approach introduces three methodological improvements over existing AI capability forecasting:

(1) Asymmetric growth curves. We use Harvey curves (Harvey, 1984) instead of standard logistic functions. Unlike the logistic curve, the Harvey function allows for an asymmetry between the initial acceleration of benchmark progress and its deceleration as scores reach saturation (Figure 2). This asymmetry is captured by a shape parameter $\alpha > 1$, which reduces the Harvey curve to a logistic function when $\alpha = 2$. This choice avoids the logistic assumption that performance should accelerate and then decelerate at the same rate. A sensitivity analysis in Appendix B confirms that the main saturation finding is robust to replacing the Harvey curve with a logistic.

(2) Hierarchical Bayesian modeling. Rather than fitting each benchmark independently, we jointly model all 98 benchmarks with shared hyperpriors over growth rates, upper asymptotes, as well as shape and noise parameters. This allows information exchange, as data-rich benchmarks inform projections for newer or more data-sparse benchmarks. The cross-validation (LOO-ELPD) model comparison confirms that joint models achieve better complexity-corrected fits (Supp. Table 2). The main saturation finding is robust to removing the hierarchical structure (Appendix B). Appendix A gives every prior, and Appendix A.5 explains what the Bayesian hierarchy does in plain terms. Models are written in PyMC (Abril-Pla et al., 2023) and sampled with the No-U-Turn Sampler (NUTS) as implemented in nutpie.

(3) Skewed likelihood at the frontier. Benchmark scores typically underestimate true model capabilities due to suboptimal prompting, scaffolding, or evaluation conditions (the "elicitation gap"). They also cannot estimate the capabilities of unreleased models. To mitigate this gap, we model observations at the top-3 frontier with a skew-normal likelihood, allowing scores to fall predominantly below a latent curve that represents the performance frontier. We verify that this modeling choice does not inflate our projections. Replacing the skew-normal with a symmetric normal likelihood yields qualitatively similar saturation conclusions (Appendix B), and moves projected saturation dates later by 2.5 months on average over the other modeling choices (Supp. Table 4).

Mathematically, for benchmark i at time t, the latent frontier capability $\mu_i(t)$ rises from the random-chance floor ℓ_i to an inferred asymptote L_i along a Harvey curve (Harvey, 1984) with growth rate k_i, inflection time τ_i and shape $\alpha_i > 1$, and the observed scores $y_i(t)$ scatter around it:

$$\mu_i(t) = \ell_i + (L_i - \ell_i)\left[1 - (1-\alpha_i)e^{-k_i(t-\tau_i)}\right]^{\frac{1}{1-\alpha_i}}, \qquad y_i(t) \sim \text{SkewNormal}\left(\mu_i(t),\, \xi_i(t),\, s_i\right). \tag{1}$$

We validate our forecasting methodology through temporal holdout, training on data before January 2025 and evaluating predictions on observed scores from January 2025 to October 2026. The joint model and the Harvey curve consistently improve the held-out continuous ranked probability score (CRPS), whereas the skew-normal likelihood predicts nearly as well as the normal one (CRPS 0.055 against 0.054 for the joint

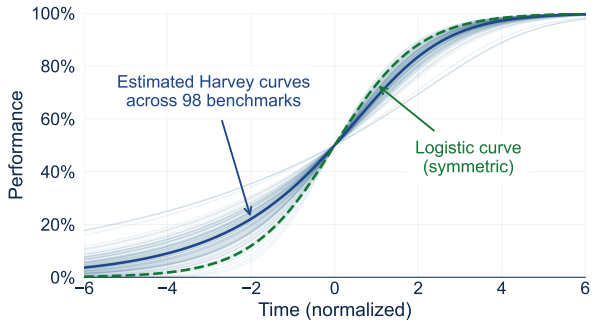


Figure 2: **Harvey curves capture asymmetric progress.** Fitted Harvey curves (blue) show more gradual acceleration compared to the logistic (green dashed). Curves are standardized to 50% at time 0 with growth rate 1. The thick blue curve is the median Harvey trajectory.

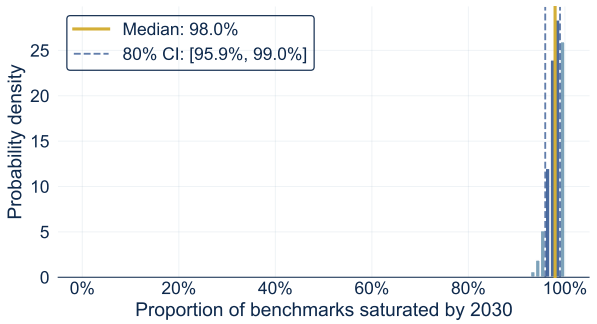


Figure 3: **Nearly all benchmarks projected to saturate by 2030.** Posterior distribution of the proportion of benchmarks reaching 95% of their score range by 2030. The median is 98.0%, indicating that saturation by 2030 is the default scenario.

Harvey variants, Supp. Table 7). We keep the skew-normal likelihood for our main model because it gives the best cross-validation fit by a wide margin (Supp. Table 2) and models the elicitation gap explicitly. On this holdout, the observations the model misses almost all fall above its 80% intervals, so its errors have been to project progress too slowly. All uncertainty intervals reported in this paper are Bayesian credible intervals (posterior quantiles). We additionally check, with conformal prediction, how much the intervals calibrated on one set of benchmarks need to be adjusted to reach their nominal coverage on benchmarks held out from calibration (Supp. Table 7). See Appendix A for model specification, Appendix B for the sensitivity analysis, and Appendix C for validation details.

2.2 Main Finding: Saturation by 2030

Our central empirical finding is that **approximately 98% of analyzed benchmarks are projected to reach saturation before 2030** (Figure 3), where saturation is defined as achieving 95% of their score range (between random chance and estimated asymptote). Benchmarks are typically considered saturated well before this point, and the smooth sigmoidal fit spreads the final percentage points over an extended period with no meaningful capability gains. Sensitivity analyses with thresholds of 90% and 99% confirm that this choice does not qualitatively affect our conclusions (Appendix B).

This finding is robust across benchmark categories, as shown in Table 1, which summarizes the twelve capability categories and projected saturation dates. All median saturation dates but three fall before 2029 (Supp. Table 8), the exceptions being FrontierMath Erdős in July 2029, DeepResearch Bench in August 2029 and CL-bench in January 2030, the only benchmark whose median falls after 2029; FrontierMath Erdős and CL-bench also carry the widest credible intervals. Even at the late end of their 80% credible intervals, 93 of the 98 benchmarks saturate before 2030. Mathematics benchmarks, whose expert baselines include a three-time IMO gold medalist at 90% on MATH and teams of mathematics students and experts solving 35% of FrontierMath problems in four and a half hours (Hendrycks et al., 2021; Epoch AI, 2025d), are mostly at or near saturation. By September 2026, frontier models scored 94% on FrontierMath Tiers 1–3 and 100% on Tier 4, and only the recent FrontierMath Erdős set, with four observations, is projected beyond 2028.

Benchmarks with direct security implications all saturate within the next few years (Figure 4). All thirteen cybersecurity benchmarks are projected to saturate by mid-2027, and all fourteen autonomous software engineering and AI R&D benchmarks by the end of 2027, suggesting that AI systems able to significantly accelerate AI research may emerge within a few years. In agentic computer use, every benchmark but DeepResearch Bench saturates by the end of 2027. Among the dual-use benchmarks, chemistry saturates earliest, by early 2026 (Supp. Figure 6), and biology by the end of 2027, with frontier models already exceeding the domain-expert baselines (38–83%) on six of the eleven biology and chemistry benchmarks that have one, and within half a point of it on a seventh (GPQA Main Chemistry). How much a saturated benchmark says about real-world capability depends on the domain. A cyber or AI R&D benchmark asks for an exploit or a code change, which is exactly what an attacker or a research team needs and could use directly. What biology and chemistry benchmarks currently test is knowledge and protocol understanding, but causing harm in

Table 1: **Projected saturation by capability category**, ordered by the latest date in each category. n is the number of benchmarks. The human baseline column gives the lowest and highest human score recorded on the category's benchmarks, from crowdworkers to expert committees: a wide range comes from these different populations, not from measurements that disagree. A dash means that no benchmark in the category has a human baseline (Appendix F lists them all). The last two columns give the earliest and latest median saturation date in the category. Supp. Table 8 lists the benchmarks one by one, and Appendix E.1 describes them.

		Human	Projected saturation	
Category	n	baseline	Earliest	Latest
Chemistry	5	22–90%	2024-08	2026-01
Cyber	13	75–96%	2025-01	2027-06
Trivia & Commonsense QA	5	80–96%	2022-10	2027-08
Biology	11	22–90%	2024-01	2027-12
Autonomous SWE	14	83%	2025-07	2027-12
Core AGI Progress	5	60–100%	2026-01	2028-01
Advanced Language and Writing	5	—	2024-11	2028-02
Domain Specific Questions	8	22–94%	2023-06	2028-06
Multimodal Understanding	8	55–100%	2025-07	2028-08
Mathematics	8	0–97%	2023-09	2029-07
Agentic Computer Use	7	72–75%	2026-05	2029-08
General Reasoning	9	53–100%	2026-01	2030-01

these fields also requires materials, a laboratory, tacit skills and a successful execution (Dev et al., 2025). We therefore do not equate their saturation with providing an operational capability to unskilled actors (though it lowers the knowledge barrier). The four remaining categories are presented in Appendix D.

2.3 Historical Context on Estimating AI Progress

Our projections should be interpreted against a recent pattern of systematic underestimation of AI progress (Kučinskas et al., 2025). This contrasts with the earlier history of AI, with periods of excessive optimism about symbolic AI and expert systems, with recurrent failed predictions that human-level AI was fifteen to twenty years away, followed by "AI winters" in the 1970s and late 1980s. However, the current wave of deep learning progress, beginning around 2012, has consistently exceeded expectations. Since 2020, milestones in language understanding, mathematical reasoning, and code generation have arrived well ahead of expert forecasts (Steinhardt, 2022). For example, AI achieving gold-medal performance at the International Mathematical Olympiad was forecast for around 2030 by both domain and non-domain experts, with only 9% probability up to 2025, the year it was achieved (Kučinskas et al., 2025).

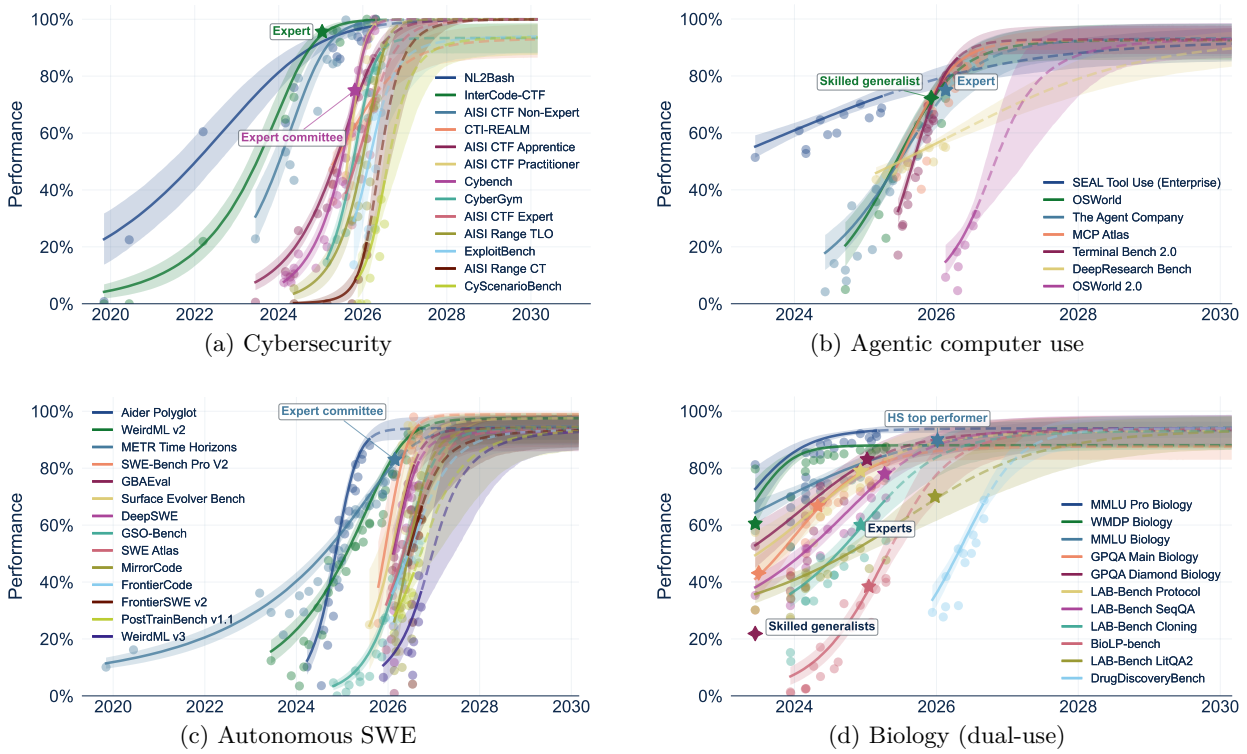
3 Implications

Our forecasts inform the time remaining before frontier systems match expert performance on most measurable cognitive tasks, but say nothing about which interventions would change that. As we do not provide evidence for the effectiveness of specific policy instruments, we keep our recommendations modest, based mainly on whether the payoff arrives inside the window.

3.1 For International Coordination

The main goal of this paper is to ensure decision-makers recognize two points. First, that GPAI capabilities are advancing rapidly, and that closely monitoring frontier progress is important for making informed policy decisions. Second, regarding control and security-critical issues, that the projected trajectory is not inevitable, as it results from investment decisions (Sevilla et al., 2024), compute infrastructure build-out (Epoch AI, 2025b), and research priorities (Erdil and Besiroglu, 2023) that remain subject to collective choice.

Given the projections in this article, the earlier investments are made to build capacity that has to exist in advance, such as monitoring of frontier capabilities, verification mechanisms and evaluation infrastructure, and the earlier multilateral discussions begin, the more options remain viable.



(a) Cybersecurity (b) Agentic computer use

(c) Autonomous SWE (d) Biology (dual-use)

Figure 4: **Security-critical capability trajectories.** Benchmarks with direct safety implications show rapid progress, with most projected to saturate by 2027–2028; chemistry, which saturated earliest, is shown in Supp. Figure 6. Starred markers indicate human baselines where available. On OSWorld, users unfamiliar with the software achieve 72% (Xie et al., 2024); on biology benchmarks, domain experts range from 38% (BioLP-bench, laboratory practitioners) to 83% (GPQA Diamond Biology, in-domain PhDs) (Dev et al., 2025; Rein et al., 2023). Graphical conventions are similar to Figure 1.

3.2 For AI Safety Research

The timeline for developing robust AI alignment techniques is short. Current approaches to alignment remain nascent. Despite some progress, they still fail to ensure a high level of safety for current AI models, let alone scaling to systems significantly exceeding human capabilities (Amodei et al., 2016; Bengio et al., 2025). Alignment is an open scientific problem, which implies that by default these systems should not be expected to remain under human control (Maier et al., 2025; Ngo et al., 2025).

Given these short timelines, one cannot rely on any single research agenda succeeding, and we are not aware of a current research direction that we would confidently describe as on track to make systems of the projected capability robustly controllable. We therefore recommend a portfolio approach:

- **Near-term agendas**: Prioritize research with potential payoffs within 3–5 years, such as technical governance agendas like robust verification mechanisms (Wasil et al., 2024; Scher and Thiergart, 2024; Peigné et al., 2026).
- **Alternative architectures**: Fund approaches that do not depend on aligning today's models, such as safe-by-design architectures or formal verification, in case alignment keeps failing in current paradigms (Dalrymple et al., 2024).
- **Differential technology development**: Favor research that makes systems safer or strengthens defenses, over research that improves AI capabilities (Sandbrink et al., 2022).

We strongly underscore, however, that safety research alone cannot guarantee good outcomes if capability development continues regardless of safety progress.

3.3 For Evaluation Practices

Forecasts, such as those presented in this article, can only be as reliable as the data upon which they are based, and as long as the underlying dynamics driving the pace of progress remain unaltered. We detail below how these two sources of uncertainty could benefit from further research attention from the benchmarking and evaluation communities.

Benchmark validity. Do benchmark improvements reflect genuine capability gains, or artifacts of optimization, contamination, or narrow task-specific learning? This concern is not novel; prior work has called for harder benchmarks targeting real-world tasks, long-horizon planning, and dynamic evaluation protocols (Mazeika et al., 2025; White et al., 2025). Our analysis already includes recently developed benchmarks designed to resist gaming, but connecting these scores to real-world impact remains understudied. Risk modeling efforts by Touzet et al. (2025); Barrett et al. (2025); Murray et al. (2025) to elicit expert knowledge in order to link capability metrics to potential harms represent a promising direction.

AI R&D acceleration. To what extent can AI systems automate AI research itself, potentially accelerating the pace of progress beyond current trends? Benchmarks measuring research capabilities (hypothesis generation, experiment design, code optimization) remain at an early stage (Kwa et al., 2025; Ihle, 2025). We encourage the development of realistic evaluations for AI R&D capabilities, as these may provide early warning of acceleration dynamics and improve medium-horizon projections.

4 Alternative Views

4.1 Benchmarks Can Be Gamed

Objection: Benchmark progress may reflect contamination (training on test data), reward-hacking on poorly designed tasks, and/or benchmark optimization (prompt engineering, fine-tuning on similar tasks, tuning scaffolds to evaluation formats) by AI companies incentivized to showcase impressive results with each new model release (Robison, 2025). Contamination in particular is a well-documented concern: Balloccu et al. (2024) show that several widely used benchmarks have appeared in training corpora. If contamination and benchmark hacking inflate scores, especially for older benchmarks with longer exposure to training pipelines, then our hierarchical model could propagate these inflated trajectories to newer, less contaminated benchmarks.

Response: We acknowledge these concerns but note several mitigating factors: (1) The saturation trends presented in this article are consistent across 98 benchmarks from four main sources (Epoch AI, Scale AI, RAND, Kaggle). (2) These benchmarks are either run internally by these organizations or sourced from carefully selected benchmark developers (Epoch AI, 2024) to limit contamination. (3) Many benchmarks saturate below 100%, which is consistent with labeling errors but inconsistent with wholesale answer memorization. (4) Newer benchmarks designed specifically to resist gaming (ARC-AGI, Remote Labor Index) show similar trajectory patterns (Chollet, 2019; Mazeika et al., 2025). (5) Some benchmarks are never publicly released, such as the UK AI Security Institute's capture-the-flag suites and cyber ranges, which rules out direct contamination. Each benchmark can be criticized, but the same pattern appears across many independently maintained benchmarks. Furthermore, fitting each benchmark independently (without hierarchical pooling) yields similar conclusions (Supp. Table 3, $\geq 91\%$ of benchmarks saturated by 2030 for all 8 model variants) with worse cross-validation fit (Supp. Table 2), suggesting that cross-benchmark information sharing does not artificially inflate projections for newer benchmarks.

4.2 Progress May Plateau

Objection: Scaling may hit diminishing returns. Data constraints, compute costs, or algorithmic limitations could slow progress before current or future benchmarks saturate. Concerns about a "data wall" (exhaustion of high-quality training data), plateauing returns from simply increasing compute (either for pre-training, post-training or inference scaling, see Ord (2025)), and the rising costs of frontier training runs all suggest that the current pace of improvement may not be sustainable. Each new scaling paradigm may buy time but may itself face diminishing returns.

Response: This is possible, and our projections carry substantial uncertainty. However, (1) predicted slowdowns in the past decade have repeatedly failed to materialize, (2) new scaling paradigms (reasoning models, inference-time compute) continue to unlock progress, as the accelerating absolute investment in AI compute and research currently compensates for diminishing returns, and (3) looking at each potential

bottleneck individually, scaling seems to be on track to continue until at least the end of the decade (Sevilla et al., 2024).

4.3 Benchmarks Do Not Measure General Intelligence

Objection: Benchmark saturation does not imply human-level intelligence. Models may reach high scores through pattern matching, heuristics or memorization, without real understanding. This concern is long-standing in AI evaluation (Chollet, 2019), given that a high score may come from optimization on a narrow task rather than from the flexible reasoning that humans transfer from one task to another. A model that solves FrontierMath problems may do so through learned solution templates rather than on-the-fly reasoning, as the impressive math skills of recent AI models seem to rely more on their extensive knowledge of the literature and of known techniques than on reasoning depth. Correspondingly, AI progress is "jagged", superhuman on some tasks yet subhuman on others. This heterogeneity could persist, preventing AGI-like systems that could reliably replace full human jobs or achieve their goals autonomously far outside of their training distribution. More generally, the field has repeatedly designated tests to measure machine intelligence, seen them passed, and moved the goalposts. Machines now pass unrestricted imitation games of the kind Turing described (Turing, 1950) and no discontinuity followed, so the benchmarks studied here may be the same mistake.

Response: Benchmark performance is indeed an imperfect proxy for general intelligence; this is why we frame our position around "measurable cognitive tasks" rather than AGI. However, (1) many recent benchmarks specifically target capabilities thought to require flexible and general reasoning (ARC Prize, 2025; Wang et al., 2025a) or the ability to handle real-world under-specified tasks (Mazeika et al., 2025), (2) as gaps are progressively closing, the breadth of saturation across diverse tasks becomes more and more difficult to explain by narrow optimization alone, and (3) if AI reaches superhuman performance on AI R&D itself, even if it is short of general intelligence, remaining gaps may close rapidly through models able to accelerate their own development.

Lastly, from a practical standpoint, systems that exceed human performance on most measurable cognitive tasks could start to have irreversible global impacts regardless of whether they possess true general intelligence. Nevertheless, we acknowledge that this is a crux and a significant source of uncertainty in predicting the consequences of benchmark saturation.

4.4 Benchmark Saturation and Superhuman Scores Do Not Entail Irreversible Global Impacts

Objection: First, benchmark saturation reflects an intrinsic property of the benchmarks' construction based on a fixed set of tasks. So even saturating all existing benchmarks and exceeding expert performance may not translate to real-world transformative capabilities. One reason is that as our analysis is limited to cognitive tasks evaluated on online benchmarks, it does not cover embodied or social capabilities which may remain bottlenecks for real-world impact. AI systems have matched or exceeded radiologist performance on medical imaging benchmarks for years, yet the demand for radiologists is greater than ever. Second, even with these wider capabilities unlocked, numerous frictions and regulatory constraints will likely delay their deployment in society. Finally, humans have always created tools that exceed their own capabilities in specific domains. Computers can process information much faster, and planes travel at superhuman speeds. AI is the latest such technology, and what its development warrants is proportionate regulation (as for the aviation industry).

Response: We agree that benchmark saturation is a necessary but not sufficient condition for irreversible real-world impact, and address each argument in turn. On the significance of benchmark saturation: the benchmarks in this article do have ceilings, so their eventual saturation is indeed expected. However, the *speed* of saturation is informative. Several benchmarks in our set were designed to resist rapid saturation (FrontierMath, Humanity's Last Exam and its HLE Diamond subset, the ARC-AGI series). That most of these benchmarks could be approaching saturation before 2028, with teams of experts on FrontierMath already well below the latest AI models, would have been considered a very optimistic prediction a few years ago. Furthermore, the benchmarks most relevant to our security argument (cybersecurity, AI R&D automation) do not require embodiment to translate into real-world impact. On deployment frictions: these are real for many beneficial applications (regulatory approval, barriers to adoption, etc.), but they are greatly reduced both for AI companies deploying AI R&D capabilities internally and for adversarial actors. Our security concerns primarily target these faster pathways. On the tools analogy: the comparison to calculators or planes understates three properties that frontier AI systems are currently acquiring and that have not been combined before. These are (a) general performance across dozens of cognitive domains, (b) autonomous

agentic operation, and (c) the potential to accelerate their own development. We believe that this combination, paired with the current absence of reliable control methods, distinguishes frontier AI from prior technologies.

More broadly, the aim of our paper is to provide evidence on the size of the window for preventive action. The consequences of saturating cybersecurity, biology, chemistry, AI R&D and remote work benchmarks are uncertain, but we note that this progress is for the most part irreversible (more so when considering the fast following of open-source models). So waiting for the gap between benchmarks and real-world impact to resolve means allowing this window for collective action to close.

4.5 What Would Change Our Mind

Several developments would substantially weaken our position. The clearest would be benchmark progress stalling for more than two years on several hard benchmarks, or evidence that current architectures hit fundamental limits on the capabilities needed to automate expert-level AI R&D. Our argument would also lose much of its force if benchmark performance were shown to diverge systematically from real-world task performance, for instance through a stall of the recent progress on the Remote Labor Index (Mazeika et al., 2025), where the best score is 21% in September 2026, or the absence of any noticeable effect of AI automation on U.S. economic growth by 2030.

5 Limitations

Extrapolation uncertainty. All forecasting involves extrapolation. Our Bayesian approach quantifies some sources of uncertainty, but it does not explicitly integrate the many known and unknown factors that may disrupt future trajectories in either direction (scaling bottlenecks, algorithmic breakthroughs, AI R&D automation, geopolitical conflicts, international agreements, etc.). An acceleration already underway would in part be absorbed by our fits, since k_i is estimated from the observed trajectory and the shape parameter lets the accelerating and decelerating phases differ in length, unlike a logistic. However, the model cannot represent a discontinuity where automated research would suddenly raise growth rates.

Interval width at long horizons. Our credible intervals understate uncertainty. Coverage of the nominal 80% interval is 75% at a horizon of 1.7 years, and lower at the earlier cutoffs, which retain only 8 to 11 benchmarks. At the 2025 and 2024 cutoffs nearly all misses fall above the interval, largely because the model underestimated the recent pace of AI progress with the rise of reasoning models; on the benchmarks that weigh most in our set it has therefore erred by projecting saturation too late rather than too early (Appendix C.4).

Elicitation gap. Benchmark scores underestimate latent model capabilities due to suboptimal prompting, scaffolding, evaluation conditions, or even sandbagging (i.e., strategic underperformance, see van der Weij et al. (2025)). Our skewed likelihood model partially addresses this issue by inferring a latent frontier capability mostly above the observed top-3 scores. Nonetheless, if substantial capabilities remain hidden due to inadequate evaluation protocols or secrecy in AI development, our projections would be conservative. The true trajectory could thus be faster than our estimates suggest.

Modeling extensions. Several extensions could improve forecast accuracy: (1) incorporating compute scaling as a predictor, linking benchmark progress to training resources (Kaplan et al., 2020; Ruan et al., 2024); (2) modeling latent capability dimensions that manifest across multiple benchmarks (Burnell et al., 2023; Kipnis et al., 2025; Ho et al., 2025; Polo et al., 2025; Pimpale et al., 2025); (3) integrating information about model architectures and training approaches. We chose a simpler model to maintain interpretability, but richer models may become feasible as more data accumulates.

Scope of analysis. Our analysis is limited to cognitive tasks evaluable through benchmarks. Embodiment, physical manipulation, and social interactions may remain critical bottlenecks for real-world impact even if cognitive benchmarks saturate. The link between benchmark saturation and transformative real-world capability remains an open question.

Benchmark evolution. Our analysis covers only existing benchmarks, and harder task sets are continuously being developed. However, in many domains, frontier models already match or exceed expert human performance, raising questions about how much headroom remains. Designing "superhuman" evaluations requires either (1) tasks where ground truth is verifiable without human judgment (e.g., mathematical proofs, code execution), or (2) aggregating across many human experts (Phan et al., 2025). Some benchmarks already approximate this standard by comparing AI to real-world human task completion. A potential next step to forecast AI progress beyond current benchmarks is to extrapolate higher-level properties of the evaluation

tasks (e.g., human-equivalent time horizon, see Kwa et al. (2025), or task difficulty level, see Zhou et al. (2026)).

6 Conclusion

We have argued that the current AI development trajectory leads to the saturation of most existing cognitive benchmarks by 2030, approaching or exceeding human expert performance across a broad range of measurable cognitive tasks. Given the state of our understanding of and control over these systems, the cost of inaction could be severe. This trajectory is not inevitable, as it results from choices that can still be influenced through coordinated action, but the window is closing rapidly.

We call on the ML community and institutional actors to engage with this possibility and act accordingly: differentially accelerating safety research, monitoring critical AI capabilities, and developing verification mechanisms as part of a global effort to provide security guarantees commensurate with the possibility of superhuman AIs in the coming years.

Reproducibility Statement

Benchmark data are publicly available from Epoch AI (Epoch AI, 2024), Scale AI (Scale AI, 2025b), RAND (Dev et al., 2025), Kaggle Open Benchmarks (Kaggle, 2026) and the Cybench leaderboard (Zhang et al., 2025). Epoch AI's benchmark data are licensed under CC BY 4.0, and the results it compiles from third parties keep their original terms; the Kaggle Open Benchmarks boards are released under Apache 2.0; scores from the RAND report are reproduced for non-commercial research under RAND's terms, and those from the Scale AI SEAL leaderboards and the Cybench leaderboard, which carry no open license, are reproduced as attributed factual results.[2] The curated scores used here (data cutoff October 1, 2026), together with the code for model fitting and visualization, are available at `https://github.com/General-Purpose-AI-Policy-Lab/benchmark-forecasting`.

Acknowledgments

We thank the anonymous reviewers for their comments, which substantially improved this paper. This work was supported by the Direction générale des relations internationales et de la stratégie (DGRIS) of the French Ministry of the Armed Forces and Veterans, through its 2026 grant programme for strategic research. The authors declare no competing interests.

References

Oriol Abril-Pla, Virgile Andreani, Colin Carroll, Larry Dong, Christopher J. Fonnesbeck, Maxim Kochurov, Ravin Kumar, Junpeng Lao, Christian C. Luhmann, Osvaldo A. Martin, Michael Osthege, Ricardo Vieira, Thomas Wiecki, and Robert Zinkov. PyMC: a modern, and comprehensive probabilistic programming framework in Python. *PeerJ Comput. Sci.*, 9:e1516, September 2023. ISSN 2376-5992. doi: 10.7717/peerj-cs.1516. URL `https://peerj.com/articles/cs-1516`.

Tom Adamczewski. OEIS Open: How many conjectures can language models turn into theorems?, August 2026. URL `https://arxiv.org/abs/2608.11941`. Version Number: 2.

Tom Adamczewski and Thomas F. Bloom. FrontierMath Erdős, September 2026. URL `https://arxiv.org/abs/2609.25050`. Version Number: 1.

Tom Adamczewski, David Owen, David Rein, Florian Brand, Giles Edkins, Allen Hart, and Daniel O'Connell. MirrorCode: AI can rebuild entire programs from behavior alone, July 2026. URL `https://arxiv.org/abs/2606.30182`. Version Number: 2.

Aider. Aider LLM Leaderboards, 2024. URL `https://aider.chat/docs/leaderboards/`.

AISI. Frontier AI Trends Report. Technical report, AI Security Institute, 2025. URL `https://www.aisi.gov.uk/frontier-ai-trends-report`.

[2]The terms and attribution of each source are listed in the repository's `NOTICE.md`. Our code and outputs are under CC BY 4.0, which does not cover the upstream scores.

Afra Feyza Akyürek, Xinming Tu, Alec Gutmanstein, Jason Qin, Divyansh Agarwal, Sofia Monasdotter, Sergey Chekhov, Brenda Hernandez Villegas, Kirill Chugunov, Judah Engel, Veronica Chatrath, Oscar Kavanagh, Geobio Boo, Ernesto Hernandez, Ying Liu, Emily Xue, Aakash Sabharwal, Daniel Yue Zhang, Zainab Doctor, Yuanhao Qu, Yunzhong He, and Sami Hassaan. DrugDiscoveryBench: Can Coding Agents Assist Early-Stage Drug Discovery? Technical report, Scale AI, 2026. URL https://static.scale.com/uploads/6691558a94899f2f65a87a75/DrugDiscoveryBench.pdf.

Dario Amodei, Chris Olah, Jacob Steinhardt, Paul Christiano, John Schulman, and Dan Mané. Concrete Problems in AI Safety, July 2016. URL http://arxiv.org/abs/1606.06565. arXiv:1606.06565 [cs].

ARC Prize. ARC-AGI-1, 2019. URL https://arcprize.org/arc-agi/1/.

ARC Prize. ARC-AGI-2, 2025. URL https://arcprize.org/arc-agi/2/.

Simone Balloccu, Patrícia Schmidtová, Mateusz Lango, and Ondřej Dušek. Leak, Cheat, Repeat: Data Contamination and Evaluation Malpractices in Closed-Source LLMs, February 2024. URL http://arxiv.org/abs/2402.03927. arXiv:2402.03927 [cs].

Chaithanya Bandi, Razvan-Gabriel Dumitru, Ben Hertzberg, Divyansh Agarwal, Geobio Boo, Tejas Polakam, Sami Hassaan, Jeff Da, HiJae Kim, Vipul Gupta, Manasi Sharma, Andrew Park, Martin Dimakis, Ernesto Gabriel Hernandez Montoya, Dan Rambado, Ivan Salazar, Rafael Cruz, MohammadHossein Rezaei, Chetan Rane, Ben Levin, Daniel Yue Zhang, Brad Kenstler, and Bing Liu. MCP-Atlas: A Large-Scale Benchmark for Tool-Use Competency with Real MCP Servers, May 2026. URL http://arxiv.org/abs/2602.00933. arXiv:2602.00933 [cs.SE].

Steve Barrett, Malcolm Murray, Otter Quarks, Matthew Smith, Jakub Kryś, Siméon Campos, Alejandro Tlaie Boria, Chloé Touzet, Sevan Hayrapet, Fred Heiding, Omer Nevo, Adam Swanda, Jair Aguirre, Asher Brass Gershovich, Eric Clay, Ryan Fetterman, Mario Fritz, Marc Juarez, Vasilios Mavroudis, and Henry Papadatos. Toward Quantitative Modeling of Cybersecurity Risks Due to AI Misuse, December 2025. URL http://arxiv.org/abs/2512.08864. arXiv:2512.08864 [cs].

Yoshua Bengio, Sören Mindermann, Daniel Privitera, Tamay Besiroglu, Rishi Bommasani, Stephen Casper, Yejin Choi, Philip Fox, Ben Garfinkel, Danielle Goldfarb, Hoda Heidari, Anson Ho, Sayash Kapoor, Leila Khalatbari, Shayne Longpre, Sam Manning, Vasilios Mavroudis, Mantas Mazeika, Julian Michael, Jessica Newman, Kwan Yee Ng, Chinasa T. Okolo, Deborah Raji, Girish Sastry, Elizabeth Seger, Theodora Skeadas, Tobin South, Emma Strubell, Florian Tramèr, Lucia Velasco, Nicole Wheeler, Daron Acemoglu, Olubayo Adekanmbi, David Dalrymple, Thomas G. Dietterich, Edward W. Felten, Pascale Fung, Pierre-Olivier Gourinchas, Fredrik Heintz, Geoffrey Hinton, Nick Jennings, Andreas Krause, Susan Leavy, Percy Liang, Teresa Ludermir, Vidushi Marda, Helen Margetts, John McDermid, Jane Munga, Arvind Narayanan, Alondra Nelson, Clara Neppel, Alice Oh, Gopal Ramchurn, Stuart Russell, Marietje Schaake, Bernhard Schölkopf, Dawn Song, Alvaro Soto, Lee Tiedrich, Gaël Varoquaux, Andrew Yao, Ya-Qin Zhang, Fahad Albalawi, Marwan Alserkal, Olubunmi Ajala, Guillaume Avrin, Christian Busch, André Carlos Ponce de Leon Ferreira de Carvalho, Bronwyn Fox, Amandeep Singh Gill, Ahmet Halit Hatip, Juha Heikkilä, Gill Jolly, Ziv Katzir, Hiroaki Kitano, Antonio Krüger, Chris Johnson, Saif M. Khan, Kyoung Mu Lee, Dominic Vincent Ligot, Oleksii Molchanovskyi, Andrea Monti, Nusu Mwamanzi, Mona Nemer, Nuria Oliver, José Ramón López Portillo, Balaraman Ravindran, Raquel Pezoa Rivera, Hammam Riza, Crystal Rugege, Ciarán Seoighe, Jerry Sheehan, Haroon Sheikh, Denise Wong, and Yi Zeng. International AI Safety Report. Technical Report arXiv:2501.17805, International AI Safety Report, January 2025. URL http://arxiv.org/abs/2501.17805. arXiv:2501.17805 [cs].

Yonatan Bisk, Rowan Zellers, Ronan Le Bras, Jianfeng Gao, and Yejin Choi. PIQA: Reasoning about Physical Commonsense in Natural Language. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 34, pages 7432–7439, April 2020. doi: 10.1609/aaai.v34i05.6239. URL https://ojs.aaai.org/index.php/AAAI/article/view/6239.

Chase Brower. Visual Physics Comprehension Test, 2025. URL https://cbrower.dev/vpct.

Ryan Burnell, Han Hao, Andrew R. A. Conway, and Jose Hernandez Orallo. Revealing the structure of language model capabilities, June 2023. URL http://arxiv.org/abs/2306.10062. arXiv:2306.10062 [cs].

Stephen Casper, Xander Davies, Claudia Shi, Thomas Krendl Gilbert, Jérémy Scheurer, Javier Rando, Rachel Freedman, Tomasz Korbak, David Lindner, Pedro Freire, Tony Wang, Samuel Marks, Charbel-Raphaël Segerie, Micah Carroll, Andi Peng, Phillip Christoffersen, Mehul Damani, Stewart Slocum, Usman Anwar, Anand Siththaranjan, Max Nadeau, Eric J. Michaud, Jacob Pfau, Dmitrii Krasheninnikov, Xin Chen, Lauro Langosco, Peter Hase, Erdem Bıyık, Anca Dragan, David Krueger, Dorsa Sadigh, and Dylan Hadfield-Menell. Open Problems and Fundamental Limitations of Reinforcement Learning from Human Feedback, September 2023. URL http://arxiv.org/abs/2307.15217. arXiv:2307.15217 [cs].

ccmdi. GeoBench, January 2026. URL https://github.com/ccmdi/geobench. original-date: 2025-03-22T13:27:48Z.

Arjun Chakraborty, Sandra Ho, Adam Cook, and Manuel Meléndez. CTI-REALM: Benchmark to Evaluate Agent Performance on Security Detection Rule Generation Capabilities, 2026. URL https://arxiv.org/abs/2603.13517. Version Number: 2.

Timothée Chauvin, Alexander Barry, Jean-Stanislas Denain, and Anson Ho. Are Mythos' cyber capabilities overhyped?, June 2026. URL https://epoch.ai/gradient-updates/are-mythos-cyber-capabilities-overhyped. Published: Epoch AI, Gradient Updates.

Evan Chen. OTIS Mock AIME 2025 Report. Technical report, 2025. URL https://web.evanchen.cc/exams/sols-OTIS-Mock-AIME-2025.pdf.

François Chollet. On the Measure of Intelligence, November 2019. URL http://arxiv.org/abs/1911.01547. arXiv:1911.01547 [cs].

Peter Clark, Isaac Cowhey, Oren Etzioni, Tushar Khot, Ashish Sabharwal, Carissa Schoenick, and Oyvind Tafjord. Think you have Solved Question Answering? Try ARC, the AI2 Reasoning Challenge, March 2018. URL http://arxiv.org/abs/1803.05457. arXiv:1803.05457 [cs].

Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. Training Verifiers to Solve Math Word Problems, November 2021. URL http://arxiv.org/abs/2110.14168. arXiv:2110.14168 [cs].

Cognition. FrontierCode Leaderboard, July 2026. URL https://cognition.com/frontiercode.

Emery Cooper, Caspar Oesterheld, Linh Chi Nguyen, Alexander Kastner, and Ethan Perez. A dataset of rated conceptual arguments, July 2026. URL https://arxiv.org/abs/2607.27499. Version Number: 1.

David "davidad" Dalrymple, Joar Skalse, Yoshua Bengio, Stuart Russell, Max Tegmark, Sanjit Seshia, Steve Omohundro, Christian Szegedy, Ben Goldhaber, Nora Ammann, Alessandro Abate, Joe Halpern, Clark Barrett, Ding Zhao, Tan Zhi-Xuan, Jeannette Wing, and Joshua Tenenbaum. Towards Guaranteed Safe AI: A Framework for Ensuring Robust and Reliable AI Systems, July 2024. URL http://arxiv.org/abs/2405.06624. arXiv:2405.06624 [cs].

Datacurve. DeepSWE, 2026. URL https://deepswe.datacurve.ai/.

Xiang Deng, Jeff Da, Edwin Pan, Yannis Yiming He, Charles Ide, Kanak Garg, Niklas Lauffer, Andrew Park, Nitin Pasari, Chetan Rane, Karmini Sampath, Maya Krishnan, Srivatsa Kundurthy, Sean Hendryx, Zifan Wang, Vijay Bharadwaj, Jeff Holm, Raja Aluri, Chen Bo Calvin Zhang, Noah Jacobson, Bing Liu, and Brad Kenstler. SWE-Bench Pro: Can AI Agents Solve Long-Horizon Software Engineering Tasks?, November 2025. URL http://arxiv.org/abs/2509.16941. arXiv:2509.16941 [cs].

Sunishchal Dev, Charles Teague, Grant Ellison, Kyle Brady, Ying-Chiang Jeffrey Lee, Sarah L. Gebauer, Henry Alexander Bradley, Dawid Maciorowski, Bria Persaud, Jordan Despanie, Barbara Del Castello, Alyssa Worland, Michael Miller, Adrian Salas, Dave Nguyen, James Liu, Jason Johnson, Andrew Sloan, Will Stonehouse, Travis Merrill, Thomas Goode, Greg McKelvey, and Ella Guest. Toward Comprehensive Benchmarking of the Biological Knowledge of Frontier Large Language Models. Technical report, RAND Corporation, 2025. URL https://www.rand.org/pubs/research_reports/RRA3797-1.html.

Shihan Dou, Ming Zhang, Zhangyue Yin, Chenhao Huang, Yujiong Shen, Junzhe Wang, Jiayi Chen, Yuchen Ni, Junjie Ye, Cheng Zhang, Huaibing Xie, Jianglu Hu, Shaolei Wang, Weichao Wang, Yanling Xiao, Yiting Liu, Zenan Xu, Zhen Guo, Pluto Zhou, Tao Gui, Zuxuan Wu, Xipeng Qiu, Qi Zhang, Xuanjing Huang, Yu-Gang Jiang, Di Wang, and Shunyu Yao. CL-bench: A Benchmark for Context Learning, 2026. URL https://arxiv.org/abs/2602.03587. Version Number: 1.

Epoch AI. AI Benchmarking Hub, November 2024. URL https://epoch.ai/benchmarks/use-this-data.

Epoch AI. CadEval, 2025a. URL https://epoch.ai/benchmarks/cad-eval.

Epoch AI. Data on Frontier AI Data Centers, September 2025b. URL https://epoch.ai/data/data-centers.

Epoch AI. Fiction.LiveBench, February 2025c. URL https://epoch.ai/benchmarks/fictionlivebench.

Epoch AI. FrontierMath, 2025d. URL https://epoch.ai/benchmarks/frontiermath.

Epoch AI. OTIS Mock AIME 2024-2025, 2025e. URL https://epoch.ai/benchmarks/otis-mock-aime-2024-2025.

Epoch AI. Earthborne Rangers (EBR-bench), July 2026a. URL https://epoch.ai/benchmarks/ebr-bench. Publication Title: Epoch AI.

Epoch AI. Furniture Assembly, September 2026b. URL https://epoch.ai/benchmarks/furniture-assembly. Publication Title: Epoch AI.

Epoch AI. Mystery Game Puzzles, July 2026c. URL https://epoch.ai/benchmarks/mystery-game-puzzles. Publication Title: Epoch AI.

Ege Erdil and Tamay Besiroglu. Algorithmic progress in computer vision, August 2023. URL http://arxiv.org/abs/2212.05153. arXiv:2212.05153 [cs].

Alexander R. Fabbri, Diego Mares, Jorge Flores, Meher Mankikar, Ernesto Hernandez, Dean Lee, Bing Liu, and Chen Xing. MultiNRC: A Challenging and Native Multilingual Reasoning Evaluation Benchmark for LLMs, July 2025. URL http://arxiv.org/abs/2507.17476. arXiv:2507.17476 [cs].

FutureSearch, Nikos I. Bosse, Jon Evans, Robert G. Gambee, Daniel Hnyk, Peter Mühlbacher, Lawrence Phillips, Dan Schwarz, and Jack Wildman. Deep Research Bench: Evaluating AI Web Research Agents, 2025. URL https://arxiv.org/abs/2506.06287. Version Number: 1.

Suhaas Garre, Emily Ritchie, Sushant Mehta, and Edwin Chen. GDP.pdf: Benchmarking Grounded Multimodal Reasoning over Professional PDF Documents, July 2026. URL https://arxiv.org/abs/2607.11192. Version Number: 3.

Advait Gosai, Tyler Vuong, Utkarsh Tyagi, Steven Li, Wenjia You, Miheer Bavare, Arda Uçar, Zhongwang Fang, Brian Jang, Bing Liu, and Yunzhong He. Audio MultiChallenge: A Multi-Turn Evaluation of Spoken Dialogue Systems on Natural Human Interaction, December 2025. URL http://arxiv.org/abs/2512.14865. arXiv:2512.14865 [cs].

Katja Grace, John Salvatier, Allan Dafoe, Baobao Zhang, and Owain Evans. When Will AI Exceed Human Performance? Evidence from AI Experts, May 2018. URL http://arxiv.org/abs/1705.08807. arXiv:1705.08807 [cs].

Katja Grace, Harlan Stewart, Julia Fabienne Sandkühler, Stephen Thomas, Ben Weinstein-Raun, Jan Brauner, and Richard C. Korzekwa. Thousands of AI Authors on the Future of AI. *Journal of Artificial Intelligence Research*, 84, October 2025. ISSN 1076-9757. doi: 10.1613/jair.1.19087. URL http://arxiv.org/abs/2401.02843. arXiv:2401.02843 [cs].

Xingang Guo, Utkarsh Tyagi, Advait Gosai, Paula Vergara, Jayeon Park, Ernesto Gabriel Hernández Montoya, Chen Bo Calvin Zhang, Bin Hu, Yunzhong He, Bing Liu, and Rakshith Sharma Srinivasa. Beyond Seeing: Evaluating Multimodal LLMs on Tool-Enabled Image Perception, Transformation, and Reasoning, October 2025. URL http://arxiv.org/abs/2510.12712. arXiv:2510.12712 [cs].

Lukas Haas, Gal Yona, Giovanni D'Antonio, Sasha Goldshtein, and Dipanjan Das. SimpleQA Verified: A Reliable Factuality Benchmark to Measure Parametric Knowledge, September 2025. URL http://arxiv.org/abs/2509.07968. arXiv:2509.07968 [cs].

A. C. Harvey. Time Series Forecasting Based on the Logistic Curve. *Journal of the Operational Research Society*, 35(7):641–646, July 1984. ISSN 0160-5682. doi: 10.1057/jors.1984.128. URL https://doi.org/10.1057/jors.1984.128. _eprint: https://doi.org/10.1057/jors.1984.128.

Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. Measuring Massive Multitask Language Understanding, 2020. URL https://arxiv.org/abs/2009.03300. Version Number: 3.

Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song, and Jacob Steinhardt. Measuring Mathematical Problem Solving With the MATH Dataset, November 2021. URL http://arxiv.org/abs/2103.03874. arXiv:2103.03874 [cs].

Yann Henon. Surface Evolver Bench, 2026. URL https://yhenon.github.io/surface-evolver-llm-eval/.

Anson Ho, Jean-Stanislas Denain, David Atanasov, Samuel Albanie, and Rohin Shah. A Rosetta Stone for AI Benchmarks, November 2025. URL http://arxiv.org/abs/2512.00193. arXiv:2512.00193 [cs].

Håvard Tveit Ihle. WeirdML, 2025. URL https://htihle.github.io/weirdml.html.

Irregular. CyScenarioBench: Evaluating LLM Cyber Capabilities Through Scenario-Based Benchmarking, 2026. URL https://irregular.com/research/cyscenariobench.

Carlos E. Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, and Karthik R. Narasimhan. SWE-bench: Can Language Models Resolve Real-world Github Issues? In *Proc. of ICLR*, pages 1–15. OpenReview.net, 2024. URL https://openreview.net/forum?id=VTF8yNQM66.

Mandar Joshi, Eunsol Choi, Daniel Weld, and Luke Zettlemoyer. TriviaQA: A Large Scale Distantly Supervised Challenge Dataset for Reading Comprehension. In *Proceedings of the 55th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1601–1611, Vancouver, Canada, 2017. Association for Computational Linguistics. doi: 10.18653/v1/P17-1147. URL http://aclweb.org/anthology/P17-1147.

Kaggle. Kaggle Benchmarks: Open Benchmarks leaderboards (MMLU, MMLU-Pro, SimpleQA Verified), 2026. URL https://www.kaggle.com/benchmarks/open-benchmarks/mmlu.

Jared Kaplan, Sam McCandlish, Tom Henighan, Tom B. Brown, Benjamin Chess, Rewon Child, Scott Gray, Alec Radford, Jeffrey Wu, and Dario Amodei. Scaling Laws for Neural Language Models, January 2020. URL http://arxiv.org/abs/2001.08361. arXiv:2001.08361 [cs].

Alex Kipnis, Konstantinos Voudouris, Luca M. Schulze Buschoff, and Eric Schulz. metabench – A Sparse Benchmark of Reasoning and Knowledge in Large Language Models, February 2025. URL http://arxiv.org/abs/2407.12844. arXiv:2407.12844 [cs].

Simas Kučinskas, Josh Rosenberg, Rebecca Ceppas de Castro, Zach Jacobs, Jordan Canedy, Philip E Tetlock, and Ezra Karger. Assessing Near-Term Accuracy in the Existential Risk Persuasion Tournament. Technical report, Forecasting Research Institute, September 2025. URL https://forecastingresearch.org/near-term-xpt-accuracy.

Thomas Kwa, Ben West, Joel Becker, Amy Deng, Katharyn Garcia, Max Hasin, Sami Jawhar, Megan Kinniment, Nate Rush, Sydney Von Arx, Ryan Bloom, Thomas Broadley, Haoxing Du, Brian Goodrich, Nikola Jurkovic, Luke Harold Miles, Seraphina Nix, Tao Lin, Neev Parikh, David Rein, Lucas Jun Koba Sato, Hjalmar Wijk, Daniel M. Ziegler, Elizabeth Barnes, and Lawrence Chan. Measuring AI Ability to Complete Long Tasks, March 2025. URL http://arxiv.org/abs/2503.14499. arXiv:2503.14499 [cs].

Jon M. Laurent, Joseph D. Janizek, Michael Ruzo, Michaela M. Hinks, Michael J. Hammerling, Siddharth Narayanan, Manvitha Ponnapati, Andrew D. White, and Samuel G. Rodriques. LAB-Bench: Measuring Capabilities of Language Models for Biology Research, July 2024. URL http://arxiv.org/abs/2407.10362. arXiv:2407.10362 [cs].

Seunghyun Lee and David Brumley. ExploitBench, 2026. URL https://exploitbench.ai/.

Nathaniel Li, Alexander Pan, Anjali Gopal, Summer Yue, Daniel Berrios, Alice Gatti, Justin D. Li, Ann-Kathrin Dombrowski, Shashwat Goel, Long Phan, Gabriel Mukobi, Nathan Helm-Burger, Rassin Lababidi, Lennart Justen, Andrew B. Liu, Michael Chen, Isabelle Barrass, Oliver Zhang, Xiaoyuan Zhu, Rishub Tamirisa, Bhrugu Bharathi, Adam Khoja, Zhenqi Zhao, Ariel Herbert-Voss, Cort B. Breuer, Samuel Marks, Oam Patel, Andy Zou, Mantas Mazeika, Zifan Wang, Palash Oswal, Weiran Lin, Adam A. Hunt, Justin Tienken-Harder, Kevin Y. Shih, Kemper Talley, John Guan, Russell Kaplan, Ian Steneker, David Campbell, Brad Jokubaitis, Alex Levinson, Jean Wang, William Qian, Kallol Krishna Karmakar, Steven Basart, Stephen Fitz, Mindy Levine, Ponnurangam Kumaraguru, Uday Tupakula, Vijay Varadharajan, Ruoyu Wang, Yan Shoshitaishvili, Jimmy Ba, Kevin M. Esvelt, Alexandr Wang, and Dan Hendrycks. The WMDP Benchmark: Measuring and Reducing Malicious Use With Unlearning, May 2024a. URL http://arxiv.org/abs/2403.03218. arXiv:2403.03218 [cs].

Qintong Li, Leyang Cui, Xueliang Zhao, Lingpeng Kong, and Wei Bi. GSM-Plus: A Comprehensive Benchmark for Evaluating the Robustness of LLMs as Mathematical Problem Solvers, July 2024b. URL http://arxiv.org/abs/2402.19255. arXiv:2402.19255 [cs].

Xi Victoria Lin, Chenglong Wang, Luke Zettlemoyer, and Michael D. Ernst. NL2Bash: A Corpus and Semantic Parser for Natural Language Interface to the Linux Operating System, 2018. URL https://arxiv.org/abs/1802.08979. Version Number: 2.

Antoine Maier, Aude Maier, and Tom David. Take Goodhart Seriously: Principled Limit on General-Purpose AI Optimization, October 2025. URL http://arxiv.org/abs/2510.02840. arXiv:2510.02840 [cs].

Mantas Mazeika, Alice Gatti, Cristina Menghini, Udari Madhushani Sehwag, Shivam Singhal, Yury Orlovskiy, Steven Basart, Manasi Sharma, Denis Peskoff, Elaine Lau, Jaehyuk Lim, Lachlan Carroll, Alice Blair, Vinaya Sivakumar, Sumana Basu, Brad Kenstler, Yuntao Ma, Julian Michael, Xiaoke Li, Oliver Ingebretsen, Aditya Mehta, Jean Mottola, John Teichmann, Kevin Yu, Zaina Shaik, Adam Khoja, Richard Ren, Jason Hausenloy, Long Phan, Ye Htet, Ankit Aich, Tahseen Rabbani, Vivswan Shah, Andriy Novykov, Felix Binder, Kirill Chugunov, Luis Ramirez, Matias Geralnik, Hernán Mesura, Dean Lee, Ed-Yeremai Hernandez

Cardona, Annette Diamond, Summer Yue, Alexandr Wang, Bing Liu, Ernesto Hernandez, and Dan Hendrycks. Remote Labor Index: Measuring AI Automation of Remote Work, October 2025. URL http://arxiv.org/abs/2510.26787. arXiv:2510.26787 [cs].

Lech Mazur. lechmazur/writing, January 2026. URL https://github.com/lechmazur/writing. original-date: 2025-01-05T08:48:04Z.

Mechanize. GBA Eval, 2026. URL https://gbaeval.com/.

Mike A. Merrill, Alexander G. Shaw, Nicholas Carlini, Boxuan Li, Harsh Raj, Ivan Bercovich, Lin Shi, Jeong Yeon Shin, Thomas Walshe, E. Kelly Buchanan, Junhong Shen, Guanghao Ye, Haowei Lin, Jason Poulos, Maoyu Wang, Marianna Nezhurina, Jenia Jitsev, Di Lu, Orfeas Menis Mastromichalakis, Zhiwei Xu, Zizhao Chen, Yue Liu, Robert Zhang, Leon Liangyu Chen, Anurag Kashyap, Jan-Lucas Uslu, Jeffrey Li, Jianbo Wu, Minghao Yan, Song Bian, Vedang Sharma, Ke Sun, Steven Dillmann, Akshay Anand, Andrew Lanpouthakoun, Bardia Koopah, Changran Hu, Etash Guha, Gabriel H. S. Dreiman, Jiacheng Zhu, Karl Krauth, Li Zhong, Niklas Muennighoff, Robert Amanfu, Shangyin Tan, Shreyas Pimpalgaonkar, Tushar Aggarwal, Xiangning Lin, Xin Lan, Xuandong Zhao, Yiqing Liang, Yuanli Wang, Zilong Wang, Changzhi Zhou, David Heineman, Hange Liu, Harsh Trivedi, John Yang, Junhong Lin, Manish Shetty, Michael Yang, Nabil Omi, Negin Raoof, Shanda Li, Terry Yue Zhuo, Wuwei Lin, Yiwei Dai, Yuxin Wang, Wenhao Chai, Shang Zhou, Dariush Wahdany, Ziyu She, Jiaming Hu, Zhikang Dong, Yuxuan Zhu, Sasha Cui, Ahson Saiyed, Arinbjörn Kolbeinsson, Jesse Hu, Christopher Michael Rytting, Ryan Marten, Yixin Wang, Alex Dimakis, Andy Konwinski, and Ludwig Schmidt. Terminal-Bench: Benchmarking Agents on Hard, Realistic Tasks in Command Line Interfaces, January 2026. URL http://arxiv.org/abs/2601.11868. arXiv:2601.11868 [cs].

Todor Mihaylov, Peter Clark, Tushar Khot, and Ashish Sabharwal. Can a Suit of Armor Conduct Electricity? A New Dataset for Open Book Question Answering, September 2018. URL http://arxiv.org/abs/1809.02789. arXiv:1809.02789 [cs].

Malcolm Murray, Steve Barrett, Henry Papadatos, Otter Quarks, Matt Smith, Alejandro Tlaie Boria, Chloé Touzet, and Siméon Campos. A Methodology for Quantitative AI Risk Modeling, December 2025. URL http://arxiv.org/abs/2512.08844. arXiv:2512.08844 [cs].

Vaskar Nath, Pranav Raja, Claire Yoon, and Sean Hendryx. ToolComp: A Multi-Tool Reasoning & Process Supervision Benchmark, 2025. URL https://arxiv.org/abs/2501.01290. Version Number: 1.

Richard Ngo, Lawrence Chan, and Sören Mindermann. The Alignment Problem from a Deep Learning Perspective, May 2025. URL http://arxiv.org/abs/2209.00626. arXiv:2209.00626 [cs].

Caspar Oesterheld, Emery Cooper, Miles Kodama, Linh Chi Nguyen, and Ethan Perez. A dataset of questions on decision-theoretic reasoning in Newcomb-like problems, 2024. URL https://arxiv.org/abs/2411.10588. Version Number: 4.

OpenAI. Learning to Reason with LLMs, September 2024. URL https://openai.com/index/learning-to-reason-with-llms.

Toby Ord. Inference Scaling Reshapes AI Governance, February 2025. URL http://arxiv.org/abs/2503.05705. arXiv:2503.05705 [cs].

Davide Paglieri, Bartłomiej Cupiał, Samuel Coward, Ulyana Piterbarg, Maciej Wolczyk, Akbir Khan, Eduardo Pignatelli, Lukasz Kuciński, Lerrel Pinto, Rob Fergus, Jakob Nicolaus Foerster, Jack Parker-Holder, and Tim Rocktäschel. BALROG: Benchmarking Agentic LLM and VLM Reasoning On Games, April 2025. URL http://arxiv.org/abs/2411.13543. arXiv:2411.13543 [cs].

Tejal Patwardhan, Rachel Dias, Elizabeth Proehl, Grace Kim, Michele Wang, Olivia Watkins, Simón Posada Fishman, Marwan Aljubeh, Phoebe Thacker, Laurance Fauconnet, Natalie S. Kim, Patrick Chao, Samuel Miserendino, Gildas Chabot, David Li, Michael Sharman, Alexandra Barr, Amelia Glaese, and Jerry Tworek. GDPval: Evaluating AI Model Performance on Real-World Economically Valuable Tasks, October 2025. URL http://arxiv.org/abs/2510.04374. arXiv:2510.04374 [cs].

Jack Payne, Jeremy Miller, and Sean Peters. Offensive Cybersecurity Time Horizons. Research Note, Lyptus Research, April 2026. URL https://lyptusresearch.org/research/offensive-cyber-time-horizons.

Pierre Peigné, Ky Nguyen, and Paul Wang. Zero knowledge verification for frontier AI training is possible, 2026. URL https://arxiv.org/abs/2606.05433. Version Number: 1.

Lukas Petersson, Axel Backlund, Axel Wennström, Hanna Petersson, Callum Sharrock, and Arash Dabiri. Blueprint-Bench: Comparing spatial intelligence of LLMs, agents and image models, 2025. URL https://arxiv.org/abs/2509.25229. Version Number: 1.

Long Phan, Alice Gatti, Ziwen Han, Nathaniel Li, Josephina Hu, Hugh Zhang, Chen Bo Calvin Zhang, Mohamed Shaaban, John Ling, Sean Shi, et al. Humanity's Last Exam, September 2025. URL http://arxiv.org/abs/2501.14249. arXiv:2501.14249 [cs].

Philip and Hemang. SimpleBench: The Text Benchmark in which Unspecialized Human Performance Exceeds that of Current Frontier Models, 2024. URL https://simple-bench.com.

Govind Pimpale, Axel Højmark, Jérémy Scheurer, and Marius Hobbhahn. Forecasting Frontier Language Model Agent Capabilities, March 2025. URL http://arxiv.org/abs/2502.15850. arXiv:2502.15850 [cs].

Felipe Maia Polo, Seamus Somerstep, Leshem Choshen, Yuekai Sun, and Mikhail Yurochkin. Sloth: scaling laws for LLM skills to predict multi-benchmark performance across families, December 2025. URL http://arxiv.org/abs/2412.06540. arXiv:2412.06540 [cs].

Proximal Labs. FrontierSWE: Benchmarking software engineering skill at the edge of human ability, 2026. URL https://www.frontierswe.com/.

Ben Rank, Hardik Bhatnagar, Ameya Prabhu, Shira Eisenberg, Karina Nguyen, Matthias Bethge, and Maksym Andriushchenko. PostTrainBench: Can LLM Agents Automate LLM Post-Training?, March 2026. URL http://arxiv.org/abs/2603.08640. arXiv:2603.08640 [cs].

Nikil Ravi, Kexing Ying, Vasilii Nesterov, Rayan Krishnan, Elif Uskuplu, Bingyu Xia, Janitha Aswedige, and Langston Nashold. ProofBench Leaderboard and Methodology, 2026. URL https://www.vals.ai/benchmarks/proof_bench. Publication Title: Vals AI.

David Rein, Betty Li Hou, Asa Cooper Stickland, Jackson Petty, Richard Yuanzhe Pang, Julien Dirani, Julian Michael, and Samuel R. Bowman. GPQA: A Graduate-Level Google-Proof Q&A Benchmark, November 2023. URL http://arxiv.org/abs/2311.12022. arXiv:2311.12022 [cs].

Kylie Robison. Meta got caught gaming AI benchmarks, April 2025. URL https://www.theverge.com/meta/645012/meta-llama-4-maverick-benchmarks-gaming.

Yaniv Romano, Evan Patterson, and Emmanuel Candes. Conformalized Quantile Regression. In H. Wallach, H. Larochelle, A. Beygelzimer, F. d' Alché-Buc, E. Fox, and R. Garnett, editors, *Advances in Neural Information Processing Systems*, volume 32, pages 3543–3553. Curran Associates, Inc., 2019. URL https://proceedings.neurips.cc/paper_files/paper/2019/file/5103c3584b063c431bd1268e9b5e76fb-Paper.pdf.

Yangjun Ruan, Chris J. Maddison, and Tatsunori Hashimoto. Observational Scaling Laws and the Predictability of Language Model Performance, October 2024. URL http://arxiv.org/abs/2405.10938. arXiv:2405.10938 [cs].

Keisuke Sakaguchi, Ronan Le Bras, Chandra Bhagavatula, and Yejin Choi. WinoGrande: An Adversarial Winograd Schema Challenge at Scale. In *The Thirty-Fourth AAAI Conference on Artificial Intelligence, AAAI 2020, The Thirty-Second Innovative Applications of Artificial Intelligence Conference, IAAI 2020, The Tenth AAAI Symposium on Educational Advances in Artificial Intelligence, EAAI 2020, New York, NY, USA, February 7-12, 2020*, pages 8732–8740. AAAI Press, 2020. URL https://aaai.org/ojs/index.php/AAAI/article/view/6399.

Jonas Sandbrink, Hamish Hobbs, Jacob Swett, Allan Dafoe, and Anders Sandberg. Differential technology development: An innovation governance consideration for navigating technology risks, September 2022. URL https://papers.ssrn.com/abstract=4213670.

Scale AI. Introducing VISTA: A Rubric-based Visual Task Assessment, 2025a. URL https://scale.com/leaderboard/visual_language_understanding.

Scale AI. SEAL LLM Leaderboards: Expert-Driven Evaluations, 2025b. URL https://scale.com/leaderboard.

Scale AI. SWE Atlas, 2026. URL https://labs.scale.com/leaderboard/sweatlas-qna. Publication Title: SEAL Leaderboard.

Aaron Scher and Lisa Thiergart. Mechanisms to Verify International Agreements About AI Development. Technical report, Machine Intelligence Research Institute, November 2024. URL https://techgov.intelligence.org/research/mechanisms-to-verify-international-agreements-about-ai-development.

Jaime Sevilla, Tamay Besiroglu, Ben Cottier, Josh You, Edu Roldan, Pablo Villalobos, and Ege Erdil. Can AI scaling continue through 2030? Technical report, Epoch AI, August 2024. URL https://epoch.ai/blog/can-ai-scaling-continue-through-2030.

Manish Shetty, Naman Jain, Jinjian Liu, Vijay Kethanaboyina, Koushik Sen, and Ion Stoica. GSO: Challenging Software Optimization Tasks for Evaluating SWE-Agents, October 2025. URL http://arxiv.org/abs/2505.23671. arXiv:2505.23671 [cs].

Ved Sirdeshmukh, Kaustubh Deshpande, Johannes Mols, Lifeng Jin, Ed-Yeremai Cardona, Dean Lee, Jeremy Kritz, Willow Primack, Summer Yue, and Chen Xing. MultiChallenge: A Realistic Multi-Turn Conversation Evaluation Benchmark Challenging to Frontier LLMs, March 2025. URL http://arxiv.org/abs/2501.17399. arXiv:2501.17399 [cs].

Rakshith S. Srinivasa, Zora Che, Chen Bo Calvin Zhang, Diego Mares, Ernesto Hernandez, Jayeon Park, Dean Lee, Guillermo Mangialardi, Charmaine Ng, Ed-Yeremai Hernandez Cardona, Anisha Gunjal, Yunzhong He, Bing Liu, and Chen Xing. TutorBench: A Benchmark To Assess Tutoring Capabilities Of Large Language Models, October 2025. URL http://arxiv.org/abs/2510.02663. arXiv:2510.02663 [cs].

Jacob Steinhardt. AI Forecasting: One Year In, July 2022. URL https://bounded-regret.ghost.io/ai-forecasting-one-year-in/.

Chloé Touzet, Henry Papadatos, Malcolm Murray, Otter Quarks, Steve Barrett, Alejandro Tlaie Boria, Elija Perrier, Matthew Smith, and Siméon Campos. The Role of Risk Modeling in Advanced AI Risk Management, December 2025. URL http://arxiv.org/abs/2512.08723. arXiv:2512.08723 [cs].

A. M. Turing. Computing Machinery and Intelligence. *Mind*, LIX(236):433–460, October 1950. ISSN 0026-4423. doi: 10.1093/mind/LIX.236.433. URL https://doi.org/10.1093/mind/LIX.236.433.

UK AI Security Institute. Our evaluation of Claude Mythos Preview's cyber capabilities, April 2026. URL https://www.aisi.gov.uk/blog/our-evaluation-of-claude-mythos-previews-cyber-capabilities.

Teun van der Weij, Felix Hofstätter, Ollie Jaffe, Samuel F. Brown, and Francis Rhys Ward. AI Sandbagging: Language Models can Strategically Underperform on Evaluations, February 2025. URL http://arxiv.org/abs/2406.07358. arXiv:2406.07358 [cs].

Aki Vehtari, Andrew Gelman, and Jonah Gabry. Practical Bayesian model evaluation using leave-one-out cross-validation and WAIC. *Stat Comput*, 27(5):1413–1432, September 2017. ISSN 1573-1375. doi: 10.1007/s11222-016-9696-4. URL https://doi.org/10.1007/s11222-016-9696-4.

Bertie Vidgen, Austin Mann, Abby Fennelly, John Wright Stanly, Lucas Rothman, Marco Burstein, Julien Benchek, David Ostrofsky, Anirudh Ravichandran, Debnil Sur, Neel Venugopal, Alannah Hsia, Isaac Robinson, Calix Huang, Olivia Varones, Daniyal Khan, Michael Haines, Austin Bridges, Jesse Boyle, Koby Twist, Zach Richards, Chirag Mahapatra, Brendan Foody, and Osvald Nitski. APEX-Agents, February 2026. URL http://arxiv.org/abs/2601.14242. arXiv:2601.14242 [cs].

Clinton J. Wang, Dean Lee, Cristina Menghini, Johannes Mols, Jack Doughty, Adam Khoja, Jayson Lynch, Sean Hendryx, Summer Yue, and Dan Hendrycks. EnigmaEval: A Benchmark of Long Multimodal Reasoning Challenges, February 2025a. URL http://arxiv.org/abs/2502.08859. arXiv:2502.08859 [cs].

Yubo Wang, Xueguang Ma, Ge Zhang, Yuansheng Ni, Abhranil Chandra, Shiguang Guo, Weiming Ren, Aaran Arulraj, Xuan He, Ziyan Jiang, Tianle Li, Max Ku, Kai Wang, Alex Zhuang, Rongqi Fan, Xiang Yue, and Wenhu Chen. MMLU-Pro: A More Robust and Challenging Multi-Task Language Understanding Benchmark. In A. Globerson, L. Mackey, D. Belgrave, A. Fan, U. Paquet, J. Tomczak, and C. Zhang, editors, *Advances in Neural Information Processing Systems*, volume 37, pages 95266–95290. Curran Associates, Inc., 2024. doi: 10.52202/079017-3018. URL https://proceedings.neurips.cc/paper_files/paper/2024/file/ad236edc564f3e3156e1b2feafb99a24-Paper-Datasets_and_Benchmarks_Track.pdf.

Zhun Wang, Tianneng Shi, Jingxuan He, Matthew Cai, Jialin Zhang, and Dawn Song. CyberGym: Evaluating AI Agents' Real-World Cybersecurity Capabilities at Scale, 2025b. URL https://arxiv.org/abs/2506.02548. Version Number: 3.

Zora Zhiruo Wang, Yijia Shao, Omar Shaikh, Daniel Fried, Graham Neubig, and Diyi Yang. How Do AI Agents Do Human Work? Comparing AI and Human Workflows Across Diverse Occupations, November 2025c. URL http://arxiv.org/abs/2510.22780. arXiv:2510.22780 [cs.AI].

Akash R. Wasil, Peter Barnett, Michael Gerovitch, Roman Hauksson, Tom Reed, and Jack William Miller. Governing dual-use technologies: Case studies of international security agreements and lessons for AI governance, September 2024. URL http://arxiv.org/abs/2409.02779. arXiv:2409.02779 [cs] version: 1.

Jason Wei, Nguyen Karina, Hyung Won Chung, Yunxin Joy Jiao, Spencer Papay, Amelia Glaese, John Schulman, and William Fedus. Measuring short-form factuality in large language models, 2024. URL https://arxiv.org/abs/2411.04368. Version Number: 1.

Kevin L. Wei, Patricia Paskov, Sunishchal Dev, Michael J. Byun, Anka Reuel, Xavier Roberts-Gaal, Rachel Calcott, Evie Coxon, and Chinmay Deshpande. Recommendations and Reporting Checklist for Rigorous & Transparent Human Baselines in Model Evaluations, 2025. URL https://arxiv.org/abs/2506.13776. Version Number: 1.

Colin White, Samuel Dooley, Manley Roberts, Arka Pal, Ben Feuer, Siddhartha Jain, Ravid Shwartz-Ziv, Neel Jain, Khalid Saifullah, Sreemanti Dey, Shubh-Agrawal, Sandeep Singh Sandha, Siddartha Naidu, Chinmay Hegde, Yann LeCun, Tom Goldstein, Willie Neiswanger, and Micah Goldblum. LiveBench: A Challenging, Contamination-Limited LLM Benchmark, April 2025. URL http://arxiv.org/abs/2406.19314. arXiv:2406.19314 [cs].

Tianbao Xie, Danyang Zhang, Jixuan Chen, Xiaochuan Li, Siheng Zhao, Ruisheng Cao, Toh Jing Hua, Zhoujun Cheng, Dongchan Shin, Fangyu Lei, Yitao Liu, Yiheng Xu, Shuyan Zhou, Silvio Savarese, Caiming Xiong, Victor Zhong, and Tao Yu. OSWorld: Benchmarking Multimodal Agents for Open-Ended Tasks in Real Computer Environments. In A. Globerson, L. Mackey, D. Belgrave, A. Fan, U. Paquet, J. Tomczak, and C. Zhang, editors, *Advances in Neural Information Processing Systems*, volume 37, pages 52040–52094. Curran Associates, Inc., 2024. doi: 10.52202/079017-1650. URL https://proceedings.neurips.cc/paper_files/paper/2024/file/5d413e48f84dc61244b6be550f1cd8f5-Paper-Datasets_and_Benchmarks_Track.pdf.

Frank F. Xu, Yufan Song, Boxuan Li, Yuxuan Tang, Kritanjali Jain, Mengxue Bao, Zora Z. Wang, Xuhui Zhou, Zhitong Guo, Murong Cao, Mingyang Yang, Hao Yang Lu, Amaad Martin, Zhe Su, Leander Maben, Raj Mehta, Wayne Chi, Lawrence Jang, Yiqing Xie, Shuyan Zhou, and Graham Neubig. TheAgentCompany: Benchmarking LLM Agents on Consequential Real World Tasks, September 2025. URL http://arxiv.org/abs/2412.14161. arXiv:2412.14161 [cs].

John Yang, Akshara Prabhakar, Karthik Narasimhan, and Shunyu Yao. InterCode: Standardizing and Benchmarking Interactive Coding with Execution Feedback, 2023. URL https://arxiv.org/abs/2306.14898. Version Number: 3.

Mengqi Yuan, Zilong Zhou, Xinzhuang Xiong, Weiming Wu, Jiayang Sun, Jiamin Song, Kaiqian Cui, Bowen Wang, Haoyuan Wu, Yitong Li, Dunjie Lu, Haikong Lu, Qi Zhen, Xinyuan Wang, Jiaqi Deng, Yuhao Yang, Cheng Chen, Boyuan Zheng, Alex Su, Xiao Yu, Hao Zou, Saaket Agashe, Xing Han Lu, Manpreet Kaur, Zhengyang Qi, Vincent Sunn Chen, Frederic Sala, Dayiheng Liu, Junyang Lin, Zhou Yu, Yu Su, Siva Reddy, Xin Eric Wang, Peng Qi, Tianbao Xie, and Tao Yu. OSWorld 2.0: Benchmarking Computer Use Agents on Long-Horizon Real-World Tasks, July 2026. URL https://arxiv.org/abs/2606.29537. Version Number: 2.

Rowan Zellers, Ari Holtzman, Yonatan Bisk, Ali Farhadi, and Yejin Choi. HellaSwag: Can a Machine Really Finish Your Sentence? In Anna Korhonen, David Traum, and Lluís Màrquez, editors, *Proc. of ACL*, pages 4791–4800, Florence, Italy, 2019. Association for Computational Linguistics. doi: 10.18653/v1/P19-1472. URL https://aclanthology.org/P19-1472.

Andy K. Zhang, Neil Perry, Riya Dulepet, Joey Ji, Celeste Menders, Justin W. Lin, Eliot Jones, Gashon Hussein, Samantha Liu, Donovan Jasper, Pura Peetathawatchai, Ari Glenn, Vikram Sivashankar, Daniel Zamoshchin, Leo Glikbarg, Derek Askaryar, Mike Yang, Teddy Zhang, Rishi Alluri, Nathan Tran, Rinnara Sangpisit, Polycarpos Yiorkadjis, Kenny Osele, Gautham Raghupathi, Dan Boneh, Daniel E. Ho, and Percy Liang. Cybench: A Framework for Evaluating Cybersecurity Capabilities and Risks of Language Models, April 2025. URL http://arxiv.org/abs/2408.08926. arXiv:2408.08926 [cs].

Lexin Zhou, Lorenzo Pacchiardi, Fernando Martínez-Plumed, Katherine M. Collins, Yael Moros-Daval, Seraphina Zhang, Qinlin Zhao, Yitian Huang, Luning Sun, Jonathan E. Prunty, Zongqian Li, Pablo Sánchez-García, Kexin Jiang-Chen, Pablo A. M. Casares, Jiyun Zu, John Burden, Behzad Mehrbakhsh, David Stillwell, Manuel Cebrian, Jindong Wang, Peter Henderson, Sherry Tongshuang Wu, Patrick C. Kyllonen, Lucy Cheke, Xing Xie, and José Hernández-Orallo. General scales unlock AI evaluation with explanatory and predictive power. *Nature*, 652(8108):58–67, April 2026. ISSN 0028-0836, 1476-4687. doi: 10.1038/s41586-026-10303-2. URL https://www.nature.com/articles/s41586-026-10303-2.

Minhui Zhu, Minyang Tian, Xiaocheng Yang, Tianci Zhou, Lifan Yuan, Penghao Zhu, Eli Chertkov, Shengyan Liu, Yufeng Du, Ziming Ji, Indranil Das, Qingzhi Chen, Junyi Cao, Jiabin Yu, Peixue Wu, Jinchen He, Yifan Su, Yikun Jiang, Yujie Zhang, Chang Liu, Ze-Min Huang, Weizhen Jia, Yunkai Wang, Farshid Jafarpour, Yong Zhao, Xinan Chen, Jessie Shelton, Aaron W. Young, John Bartolotta, Wenchao Xu, Yue Sun, Anjun Chu, Victor Colussi, Chris Akers, Nathan Brooks, Wenbo Fu, Jinchao Zhao, Marvin Qi, Anqi Mu, Yubo Yang, Allen Zang, Yang Lyu, Peizhi Mai, Christopher Wilson, Xuefei Guo, Juntai Zhou, Daniel Inafuku, Chi Xue, Luyu Gao, Ze Yang, Yaïr Hein, Yonatan Kahn, Kevin Zhou, Di Luo, John Drew Wilson,

Jarrod T. Reilly, Dmytro Bandak, Ofir Press, Liang Yang, Xueying Wang, Hao Tong, Nicolas Chia, Eliu Huerta, and Hao Peng. Probing the Critical Point (CritPt) of AI Reasoning: a Frontier Physics Research Benchmark, 2025. URL https://arxiv.org/abs/2509.26574. Version Number: 4.

A Methodological Details

A.1 Growth Curve Specification

We write $y_i(t)$ for the score observed on benchmark i at time t, counted in days since the benchmark's first observation. The latent frontier performance follows a shifted sigmoid:

$$\mu_i(t) = \ell_i + (L_i - \ell_i) \cdot \sigma_i(t),$$

Here $\ell_i \in [0, 1]$ is the lower bound of benchmark i, its random-chance score. We take it from the answer format, from a published chance baseline or from the scoring rule, and set it to 0 for free-response, agentic and rubric-graded benchmarks, where guessing scores nothing (Appendix E.2). $L_i \in [L_{\min}, 1]$ is the upper asymptote, which the model infers, with $L_{\min} = 0.75$ (see below).

Harvey function. The Harvey curve generalizes the logistic with a shape parameter $\alpha_i > 1$:

$$\sigma_i^{\text{Harvey}}(t) = [1 - (1 - \alpha_i) \exp(-k_i(t - \tau_i))]^{\frac{1}{1-\alpha_i}},$$

where $k_i > 0$ is the growth rate and τ_i the inflection time. The larger α_i, the more gradual the acceleration and the faster the final approach to the asymptote. With $\alpha_i = 2$, the Harvey function is the standard logistic.

Logistic function. For comparison, we also fit the standard logistic:

$$\sigma_i^{\text{Logistic}}(t) = \frac{1}{1 + \exp(-k_i(t - \tau_i))}.$$

A.2 Observation Model

Observations follow a skew-normal distribution with heteroskedastic noise:

$$y_i(t) \sim \text{SkewNormal}(\mu_i(t), \xi_i(t), s_i),$$

where $s_i \leq 0$ is a skewness parameter allowing scores to fall predominantly below the latent curve.

The noise scale $\xi_i(t)$ is heteroskedastic and Beta-shaped over $[\ell_i, L_i]$:

$$\xi_i(t) = \xi_0 + \xi_i^{\text{base}} \cdot \frac{\sqrt{(\mu_i(t) - \ell_i)(L_i - \mu_i(t))}}{(L_i - \ell_i)/2},$$

with $\xi_0 = 0.01$ fixed. This yields maximal variance at mid-range and minimal variance near the bounds.

A.3 Hierarchical Structure

The joint model places shared hyperpriors across benchmarks:

Upper asymptote. $L_i = L_{\min} + (1 - L_{\min}) \cdot L_i^{\text{raw}}$, with $L_{\min} = 0.75$ and $L_i^{\text{raw}} \sim \text{Beta}(\mu_L, \sigma_L)$, and hyperpriors $\mu_L \sim \text{Beta}\left(\text{mean} = \frac{0.96 - L_{\min}}{1 - L_{\min}}, \text{sd} = \frac{0.02}{1 - L_{\min}}\right)$ and $\sigma_L \sim \text{Half-}\mathcal{N}\left(\text{sd} = \frac{0.02}{1 - L_{\min}}\right)$, truncated at $\sigma_L \leq (1 - \mu_L)\sqrt{\mu_L/(2 - \mu_L)}$ so that the second shape parameter of the Beta stays at or above 1. Below that value the density of L_i diverges at 1, which would put a spike of prior and posterior mass on perfect scores. On the scale of L_i, the hyperprior on the population mean $L_{\min} + (1 - L_{\min})\,\mu_L$ is thus centered on 0.96 with standard deviation 0.02. Two constraints come from the data. First, a frontier cannot plateau below what has already been reached, so the Beta of each benchmark is cut off below its best observed performance, human baseline or model score (an interval transform on the draw, so that μ_L and σ_L keep describing the asymptotes themselves). Second, L_i is pinned at 1 for the 13 benchmarks whose tasks are all known to be solvable (ARC-AGI, ARC-AGI-2, EBR-bench, VPCT, Cybench, the two FrontierMath v2 sets and the six UK AI Security Institute capture-the-flag suites and cyber ranges), and the floor pins it at 1 for the five on which a score of 100% has already been recorded.

Growth rate. $k_i \sim \mathcal{G}(k_\mu, k_\sigma)$, with $k_\mu \sim \mathcal{G}(\text{mean} = 0.005, \text{sd} = 0.002)$ and $k_\sigma \sim \text{Half-}\mathcal{N}(\text{sd} = 0.005)$.

Inflection time. $\tau_i \sim \text{Gumbel}(\hat{\tau}_i, \beta)$, where $\hat{\tau}_i$ is the empirical midpoint of benchmark i's observed time range and $\beta = 730$ days (2 years). The right-skewed Gumbel prior reflects that for unsaturated benchmarks, the inflection point likely lies beyond observed data.

Noise. $\xi_i^{\text{base}} \sim \mathcal{G}(\xi_\mu^{\text{base}}, \xi_\sigma^{\text{base}})$, with hyperpriors $\xi_\mu^{\text{base}} \sim \mathcal{G}(\text{mean} = 0.05 + \frac{N}{50}, \text{sd} = 0.02)$ and $\xi_\sigma^{\text{base}} \sim \text{Half-}\mathcal{N}(\text{sd} = 0.05)$, where N corresponds to the top-N frontier of scores considered ($N = 3$ in this paper).

Skewness. $s_i \sim \text{Truncated-}\mathcal{N}(s_\mu, s_\sigma, -\infty, 0)$, with hyperpriors $s_\mu \sim \mathcal{N}(\text{mean} = -2 - \frac{N}{2}, \text{sd} = 0.5)$ and $s_\sigma \sim \text{Half-}\mathcal{N}(\text{sd} = 1.0)$.

Harvey shape. $\alpha_i = 1 + \alpha_i^{\text{raw}}$, with $\alpha_i^{\text{raw}} \sim \mathcal{G}(\alpha_\mu^{\text{raw}}, \alpha_\sigma^{\text{raw}})$, ensuring $\alpha_i > 1$, and hyperpriors $\alpha_\mu^{\text{raw}} \sim \mathcal{G}(\text{mean} = 1.5, \text{sd} = 0.5)$ and $\alpha_\sigma^{\text{raw}} \sim \text{Half-}\mathcal{N}(\text{sd} = 0.5)$.

A.4 Inference

Each model family enters a benchmark once per release day, with its best score across sizes and settings, and we track the top-3 frontier scores per benchmark at each point in time to reduce noise from suboptimal evaluations. Models are written in PyMC (Abril-Pla et al., 2023) and sampled with the No-U-Turn Sampler as compiled by nutpie (4 chains of 2000 posterior samples after 1000 warmup iterations at target acceptance 0.9; one draw in four is kept, so 2000 posterior draws per fit). In the independent variants, each benchmark's hyperparameters have a single member and stay at their prior, so we integrate them out by quadrature rather than sample them, and use a target acceptance of 0.95. Temporal-holdout fits (Appendix C) include fewer benchmarks, which leaves the hyperpriors less constrained, and use a target acceptance of 0.99, 2000 warmup iterations and 5000 posterior samples per chain for every variant, thinned in the same way to 5000 posterior draws per fit. Convergence is assessed via $\hat{R}$ diagnostics and effective sample size. The main fit has $\hat{R} \leq 1.01$ for every parameter and 11 divergent transitions out of 8000 sampled draws, and the other full-data variants $\hat{R} \leq 1.05$ with under 2% of divergent transitions. On a laptop CPU (Apple M4, 16 GB of memory), the 16 fits of the main analysis (8 variants on the full data and on the 2025 holdout) take about one hour, and the additional retrodiction and prior-sensitivity fits about 40 more minutes.

A.5 What the Hierarchy Does

The joint model estimates how growth rates, asymptotes, shapes and noise levels are distributed across the 98 benchmarks, and draws the parameters of each benchmark toward that distribution, especially when they have little data. For instance, rather than fixing by hand the parameters of a Beta prior on the upper asymptote, we give these parameters a prior of their own and infer the Beta distribution of asymptotes from the data.

This is not an assumption that benchmarks behave alike. The population distribution has an estimated spread, and all the parameters are still estimated for each benchmark. Nevertheless, pooling propagates error as well as information, which is why we report the independent fits throughout, and check in Supp. Table 4 that removing the hierarchy moves saturation dates later by only 2.4 months on average (median 0.7 months).

B Sensitivity Analysis

We assess the robustness of our main finding to all modeling choices by evaluating a full grid of 8 model variants: sigmoid family (Harvey vs. logistic) × hierarchical structure (joint vs. independent) × likelihood (skew-normal vs. normal). Independent fits use identical model structure but estimate all parameters separately for each benchmark without shared hyperpriors. Each variant is evaluated on model comparison via LOO-ELPD and saturation projections.

B.1 Model Comparison (LOO-ELPD)

We compare all 8 variants using leave-one-out cross-validation (LOO; Vehtari et al., 2017). LOO-ELPD estimates each model's predictive accuracy by measuring how well it predicts each observation when that observation is held out, implicitly penalizing model complexity without requiring explicit parameter counting.

Supp. Table 2 reports the expected log pointwise predictive density (ELPD; higher is better), the difference relative to the best model (ΔELPD; negative means worse), and the Bayesian stacking weight (the optimal weight for combining models into a predictive ensemble). Each of the three choices improves fit in every pairing with the other two: the Harvey curve over the logistic, the skew-normal over the normal likelihood, and the joint over the independent structure. The best logistic variant trails the main model by 91.5 ± 12.2, more than seven standard errors, and stacking assigns essentially all of its weight to the main model. Depending on the variant, Pareto-smoothed importance sampling flags 43 to 73 of the 1,621 observations as influential ($\hat{k} > 0.7$). The ELPD values are therefore approximate, but the differences between variants remain several times larger than their standard errors.

Table 2: LOO-ELPD model comparison across all 8 variants (Vehtari et al., 2017). Higher ELPD indicates better predictive fit. ΔELPD is the difference relative to the best model; ΔSE is its standard error. Weight is the Bayesian stacking weight.

Model variant	ELPD	ΔELPD	ΔSE	Weight
Harvey joint (skew)	2246.8	0	—	1.00
Harvey indep. (skew)	2190.6	−56.2	7.4	0
Harvey joint (normal)	2165.8	−81.1	12.2	≈0
Logistic joint (skew)	2155.3	−91.5	12.2	≈0
Logistic indep. (skew)	2123.2	−123.6	14.4	≈0
Harvey indep. (normal)	2121.7	−125.1	12.3	0
Logistic joint (normal)	2094.0	−152.9	15.5	0
Logistic indep. (normal)	2072.2	−174.7	15.1	0

B.2 Sensitivity to Saturation Threshold

Supp. Table 3 reports the proportion of benchmarks projected to reach saturation by 2030 across all 8 model variants and three threshold definitions (90%, 95%, 99% of score range, from random chance to their estimated upper asymptote L_i). The qualitative finding is robust to all modeling alternatives (≥91% of benchmarks saturated by 2030 at the 95% threshold). Hierarchical (joint) models produce higher saturation estimates, as expected from information sharing across benchmarks. Saturation estimates are lower with the 99% threshold, although, even under this strict definition, 92% of benchmarks are estimated to be saturated by 2030 with the best model, and ≥82% for all model variants.

Table 3: Sensitivity of saturation projections across all 8 model variants and three threshold definitions (90%, 95% and 99% of the score range, from the lower bound ℓ_i to the estimated asymptote L_i).

	Threshold: 90%		Threshold: 95%		Threshold: 99%	
Model variant	Median	80% CI	Median	80% CI	Median	80% CI
Harvey joint (skew)	99.0%	[96.9, 100]	98.0%	[95.9, 99.0]	91.8%	[88.8, 94.9]
Harvey joint (normal)	98.0%	[95.9, 100]	95.9%	[93.9, 98.0]	87.8%	[84.7, 90.8]
Harvey indep. (skew)	96.9%	[95.9, 99.0]	94.9%	[92.9, 96.9]	86.7%	[83.7, 89.8]
Harvey indep. (normal)	95.9%	[93.8, 98.0]	92.9%	[89.8, 94.9]	84.7%	[81.6, 87.8]
Logistic joint (skew)	98.0%	[96.9, 99.0]	95.9%	[93.9, 98.0]	85.7%	[83.7, 88.8]
Logistic joint (normal)	98.0%	[95.9, 99.0]	94.9%	[92.9, 96.9]	85.7%	[82.7, 88.8]
Logistic indep. (skew)	96.9%	[94.9, 99.0]	94.9%	[92.9, 96.9]	84.7%	[81.6, 86.7]
Logistic indep. (normal)	94.9%	[92.9, 96.9]	91.8%	[89.8, 94.9]	82.7%	[79.6, 85.7]

B.3 Effect of Modeling Choices on Saturation Dates

We invert each fitted curve analytically to obtain, for every benchmark and posterior draw, the date at which it reaches 95% of its score range, and compare those dates across variants (Supp. Table 4, each choice averaged over the other two). All three effects are small against the four-year horizon of our claim.

B.4 Sensitivity to the Prior on the Upper Asymptote

Because saturation is defined relative to L_i, an informative prior on it could influence the result, and ours is informative, with a floor at $L_{\min} = 0.75$ and a population mean centered on 0.96 with standard deviation

Table 4: **Effect of each modeling choice on projected saturation dates**, in months (positive means later saturation). Each row averages the contrast over the two remaining factors and over all benchmarks, using the posterior median saturation date per benchmark. Per-benchmark dates are reported in Supp. Table 8.

Modeling choice	Mean shift	Median shift	10–90% range
Skew-normal → normal likelihood	+2.5 months	+2.3 months	[-0.0, +8.2]
Joint → independent	+2.4 months	+0.7 months	[-0.7, +4.9]
Harvey → logistic	+2.4 months	+1.0 month	[-0.4, +5.3]

0.02. Supp. Table 5 refits with the floor at 0.50, with the standard deviation widened to 0.05, and with the floor at 0.50 and the standard deviation widened further to 0.10.

None of the three changes moves the headline proportion, which stays at 98.0% (80% CI: [95.9, 99.0]). Lowering the floor changes little, since no benchmark's data supports an asymptote near 0.50, and the asymptotes are in any case held above the best performance already observed (Appendix A). Widening the prior lets the asymptote estimates spread further down, from a 10th percentile of 0.928 to 0.915 with the lower floor and the widest prior, and a lower asymptote leaves a shorter range to cover before reaching 95% of it, which makes saturation easier rather than harder. Our prior is therefore not what drives the conclusion.

Table 5: **Sensitivity to the prior on the upper asymptote.** $L_{\min}$ is the hard floor of the asymptote and sd, on the scale of the asymptote, is both the standard deviation of the hyperprior on its population mean and the scale of the half-normal prior on its spread. Relaxing the prior leaves the projected proportion saturated by 2030 unchanged. The last column reports the minimum effective sample size and the number of divergent transitions; the variant that samples least well is Low floor ($L_{\min}$=0.50).

Prior on L	Sat. by 2030	80% CI	Median L	L 10–90%	ESS / div.
Main ($L_{\min}$=0.75, sd 0.02)	98.0%	[95.9, 99.0]	0.948	[0.928, 1.000]	820 / 11
Low floor ($L_{\min}$=0.50)	98.0%	[95.9, 99.0]	0.949	[0.926, 1.000]	303 / 45
Weak prior (sd 0.05)	98.0%	[95.9, 99.0]	0.942	[0.919, 1.000]	385 / 24
Low floor + weaker prior (sd 0.10)	98.0%	[95.9, 99.0]	0.942	[0.915, 1.000]	844 / 23

C Retrodiction Analysis

We validate our forecasting methodology using temporal holdout, training on data before a cutoff date and evaluating predictions on subsequently observed scores.

C.1 Protocol

We set the cutoff date to January 1, 2025. Models are trained on all benchmark data prior to this date, then asked to predict scores observed on models released between January 1, 2025 and October 1, 2026. Benchmarks with fewer than five pre-cutoff frontier observations are excluded, since a sigmoid can hardly be identified from fewer points, which leaves 49 of the 98 benchmarks and 417 held-out observations. We compare all model variants. Appendix C.4 repeats the exercise at earlier cutoffs to probe longer horizons.

C.2 Evaluation Metrics

Bayesian Predictive Calibration. We assess whether predicted credible intervals achieve their nominal coverage. For each confidence level $p \in [0.01, 0.99]$, we compute the fraction of test observations falling within the p-credible interval of the posterior predictive distribution. A well-calibrated model shows observed coverage matching expected coverage.

CRPS. The Continuous Ranked Probability Score measures the quality of probabilistic forecasts, penalizing both miscalibration and lack of sharpness. It generalizes the Brier score to continuous predictions. Lower CRPS indicates better predictions.

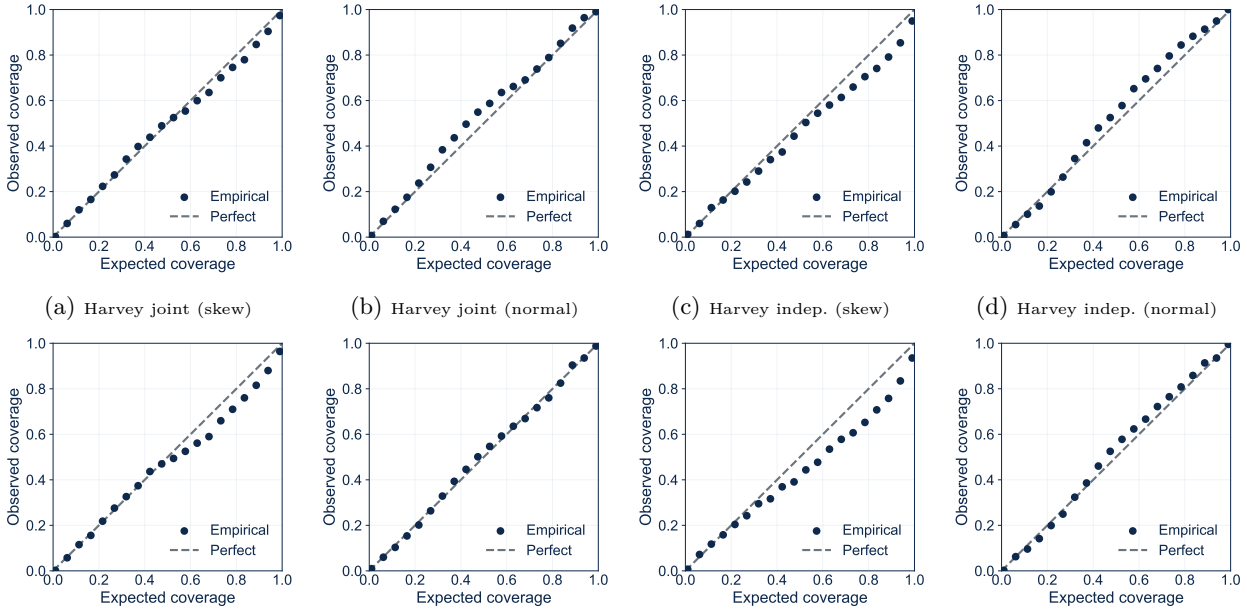
RMSE. Root Mean Squared Error between the posterior mean prediction and observed scores.

Table 6: **Retrodiction accuracy as a function of forecast horizon** for the main model (Harvey joint, skew-normal). Each row trains on frontier scores released before the cutoff and predicts every score observed between the cutoff and October 2026. The retrospective filter keeps only benchmarks with at least 5 pre-cutoff frontier observations, which is why the earlier cutoffs cover far fewer benchmarks: at the 2023 cutoff only the commonsense and early question-answering sets and GSM8K existed. Nominal coverage is 80%; the last two columns count the held-out observations falling above the upper bound of the 80% credible interval and those falling below its lower bound, so a large asymmetry over many observations means the model erred in one direction.

Train cutoff	Horizon	Benchmarks	n_{train}	n_{test}	CRPS	RMSE	80% cov.	Above	Below
2023-01	3.7 yr	8	57	88	0.043	0.070	70.5%	2	24
2024-01	2.7 yr	11	114	131	0.088	0.183	64.1%	45	2
2025-01	1.7 yr	49	561	417	0.055	0.109	75.3%	100	3

C.3 Retrodiction Results

Supp. Figure 5 showcases model calibration. Full retrodiction metrics (CRPS, RMSE) for all 8 model variants are reported in Supp. Table 7.



(a) Harvey joint (skew) (b) Harvey joint (normal) (c) Harvey indep. (skew) (d) Harvey indep. (normal)

(e) Logistic joint (skew) (f) Logistic joint (normal) (g) Logistic indep. (skew) (h) Logistic indep. (normal)

Figure 5: Calibration curves for all 8 model variants. Well-calibrated models produce points close to the diagonal. Normal-likelihood variants come closest to nominal coverage on this holdout; skew-normal variants have intervals that are somewhat too narrow.

All eight model variants achieve strong out-of-sample predictive performance (CRPS ≤ 0.062), with differences between variants that are small against the score scale (Supp. Table 7). The joint model and the Harvey curve improve CRPS in every pairing with the other choices. The skew-normal likelihood predicts nearly as well as the normal one (CRPS 0.055 against 0.054 for the joint Harvey variants), and the normal one is closer to nominal coverage (80.3% against 75.3% for the nominal 80% interval, Supp. Figure 5). For the main model, 24.0% of the held-out observations fall above the 80% interval and 0.7% below it, so the misses are one-sided and the model underpredicted the pace of progress over 2025–2026.

We select the Harvey hierarchical model with skew-normal likelihood for our main analysis. The hierarchical structure enables information sharing across benchmarks without degrading performance, which may particularly benefit predictions for data-sparse benchmarks. Although the Harvey model is more complex than the logistic, Bayesian inference naturally regularizes through prior specifications, avoiding the overfitting concerns

that arise with complex models in classical machine learning. We therefore choose the model that aligns with our theoretical understanding of capability progress and that is confirmed as optimal by formal model comparison (Supp. Table 2). The main model predicts the holdout nearly as well as the normal-likelihood variant, which leads to the same conclusion (95.9% of benchmarks saturated by 2030, Supp. Table 3) and whose saturation dates are later than those of the main model by 4.1 months on average (2.5 months for the skew-normal to normal contrast averaged over the other modeling choices, Supp. Table 4).

C.4 Forecast Horizon

The holdout above spans 21 months, whereas our claim concerns 2030. We therefore push the training cutoff back to January 2024 and 2023, giving horizons up to 3.7 years (Supp. Table 6). This exercise necessarily reduces the number of benchmarks covered, since the retrospective filter keeps only benchmarks with at least five pre-cutoff observations, so 8 survive at the 2023 cutoff (commonsense and early question-answering sets, plus GSM8K) and 11 at 2024; at the 2022 cutoff only three would, too few to report. All holdout fits converge: $\hat{R}$ is at most 1.01 for the main model at every cutoff and 1.04 across the eight variants, with under 1% of divergent transitions.

Coverage of the nominal 80% interval is 75% at the 2025 cutoff, 64% at 2024 and 70% at 2023; with so few data the coverage estimates at the earlier cutoffs are themselves imprecise. The last two columns of Supp. Table 6 show the direction of the misses. At the 2025 and 2024 cutoffs nearly all fall above the upper bound (100 observations above against 3 below, and 45 against 2), so the model projected saturation too late. The 2023 cutoff errs the other way on 24 of its 88 observations (2 above), on eight early benchmarks whose final points of progress came more slowly than projected. Our intervals are therefore too narrow at multi-year horizons, and on the benchmarks that weigh most in our set they have erred by projecting saturation too late.

C.5 Validation: Conformal Prediction Intervals

As a model-free validation of our uncertainty quantification, we apply Conformalized Quantile Regression (CQR; Romano et al., 2019) to the temporal holdout predictions for all 8 model variants. CQR is a distribution-free method that adjusts Bayesian credible intervals using held-out calibration data. It computes a correction term Q such that the adjusted intervals $[\text{lower} - Q, \text{upper} + Q]$ achieve guaranteed coverage under exchangeability. We assign whole benchmarks at random to the calibration and test halves, repeating over 100 assignments. A small $|Q|$ indicates that the Bayesian intervals are already well-calibrated and require little correction. Supp. Table 7 shows that corrections are small across all variants, at most 0.023 on the 0–1 score scale. They are positive for the skew-normal variants, whose intervals CQR widens, and close to zero or slightly negative for the normal ones, and after adjustment every variant lands between 76% and 81% coverage.

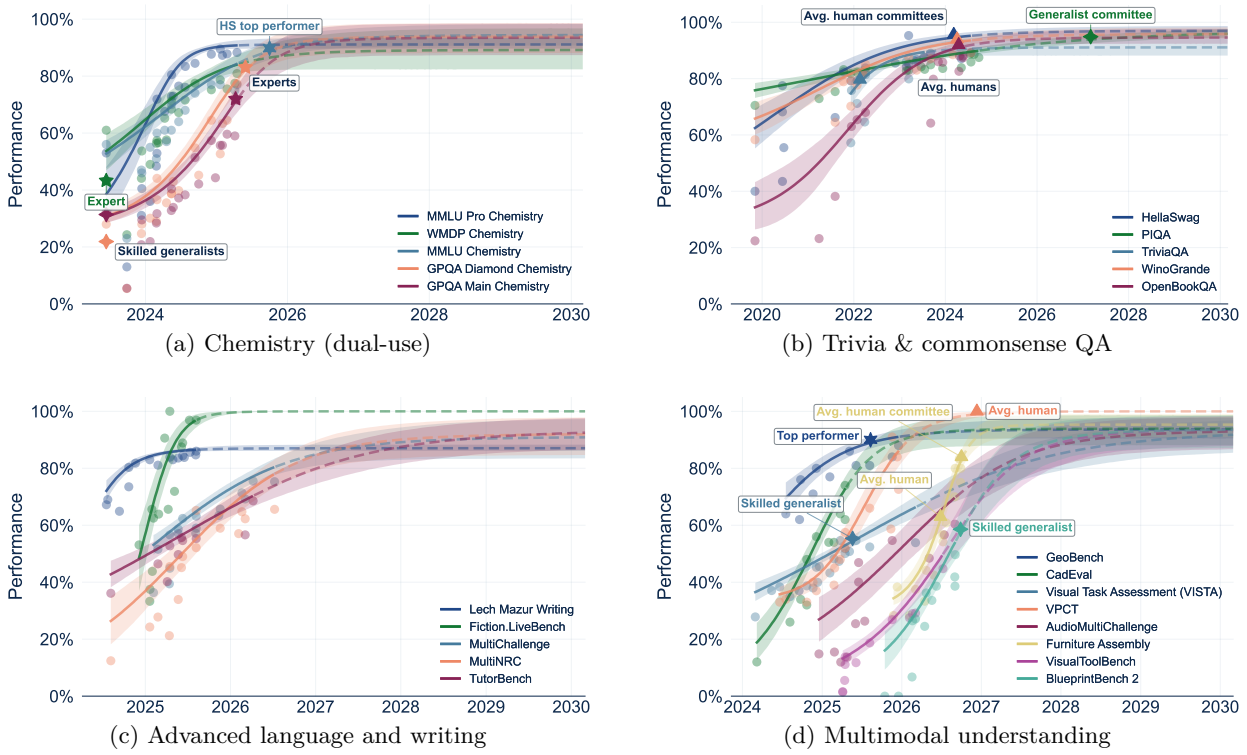
Table 7: **Held-out accuracy and distribution-free calibration.** CRPS and RMSE are out-of-sample scores on the 2025 holdout (lower is better); bold marks the variants a paired comparison cannot separate from the best. Q is the Conformalized Quantile Regression adjustment (Romano et al., 2019) that would be needed for a finite-sample coverage guarantee. Bayesian coverage is the fraction of all held-out observations inside the nominal 80% credible interval, and CQR cov. the coverage of the adjusted intervals. Calibration and test halves are formed by assigning whole benchmarks at random, so the quantity measured is whether intervals calibrated on some benchmarks cover on others; Q and CQR cov. are medians over 100 random assignments.

Model variant	CRPS	RMSE	Q	Bayes. cov.	CQR cov.
Harvey joint (skew)	0.055	0.109	+0.010	75.3%	77.2%
Harvey joint (normal)	**0.054**	**0.105**	−0.003	80.3%	76.4%
Harvey indep. (skew)	0.059	0.112	+0.015	72.2%	77.0%
Harvey indep. (normal)	0.057	**0.106**	−0.013	85.9%	80.7%
Logistic joint (skew)	0.057	0.111	+0.013	72.7%	76.7%
Logistic joint (normal)	0.057	0.110	+0.000	78.4%	77.7%
Logistic indep. (skew)	0.062	0.118	+0.023	67.6%	76.4%
Logistic indep. (normal)	0.058	**0.107**	−0.007	82.3%	78.9%

D Additional Results

D.1 Additional Categories

Supp. Figure 6 shows trajectories for the four remaining categories. The dual-use chemistry benchmarks saturated first among security-relevant categories, all five by early 2026, with frontier models well above the PhD-level baseline on WMDP Chemistry (81% against 43%) and within six points of it on the two GPQA chemistry sets (Dev et al., 2025). Trivia and commonsense benchmarks saturated earliest overall (TriviaQA in 2022, HellaSwag in 2023, WinoGrande in 2024), consistent with these tasks representing lower cognitive complexity; PIQA is the exception, with a best score of 89% so far and saturation projected in 2027. Language and writing, and multimodal benchmarks show intermediate trajectories, with saturation projected by the end of 2028.



(a) Chemistry (dual-use)

(b) Trivia & commonsense QA

(c) Advanced language and writing

(d) Multimodal understanding

Figure 6: **Additional benchmark trajectories.** Chemistry benchmarks are already saturated; trivia and commonsense benchmarks are at or near saturation; language and multimodal understanding show trajectories consistent with other categories. Graphical conventions are similar to Figure 1.

D.2 Posterior Hyperparameters

Supp. Figure 7 shows the population-level hyperparameters of the main model. The growth rate is centered on $k_\mu = 0.0084$ per day, roughly 4 percentage points per month at the steepest part of a trajectory. The shape parameter confirms the asymmetry that motivates the Harvey curve, with a median of 3.36 (80% CI: $[2.99, 3.73]$) against the value of 2 at which the curve reduces to a logistic. The asymptote population mean, $L_{\min} + (1 - L_{\min})\,\mu_L$, is 0.941 (80% CI: $[0.929, 0.951]$), and the skewness is clearly negative at -2.09 (80% CI: $[-2.48, -1.74]$).

D.3 Upper Asymptote by Benchmark

Supp. Figure 8 shows the benchmark-specific asymptotes compared to the best score observed so far. Since an asymptote cannot fall below the best performance already observed (Appendix A), the lowest estimates sit just above their benchmark's best score: 0.865 on GPQA Main Biology (best score 0.821) and 0.866 on Lech

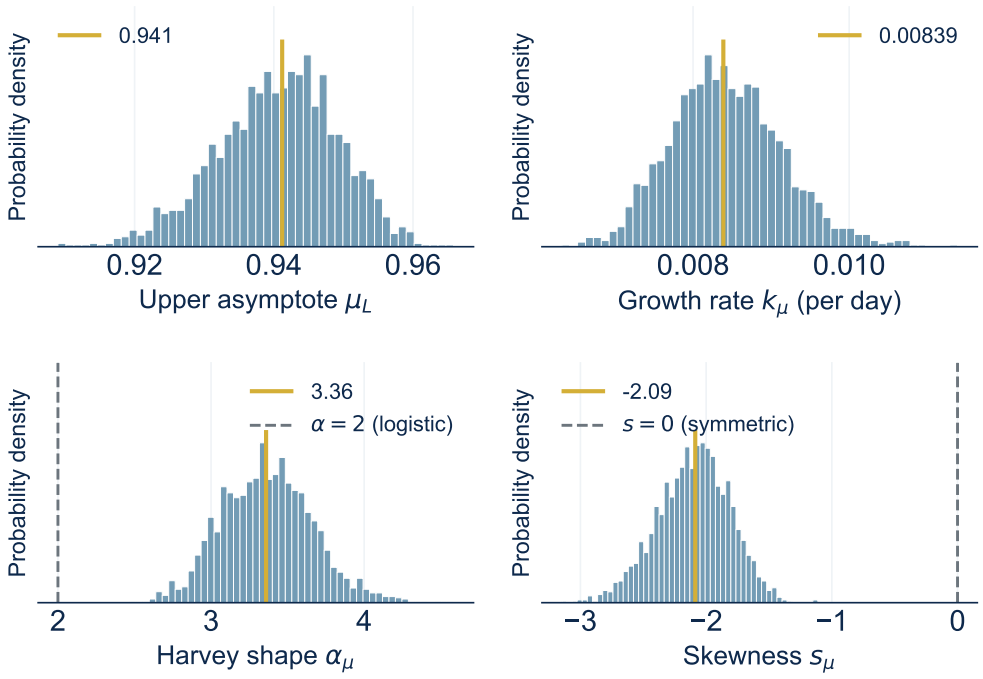


Figure 7: **Hierarchical model hyperparameters.** Posterior distributions for the joint Harvey model hyperparameters, controlling the population-level distribution of upper asymptotes ($L_{\min} + (1 - L_{\min})\,\mu_L$), growth rates ($k_\mu$), Harvey shapes ($\alpha_\mu = 1 + \alpha_\mu^{\text{raw}}$), and skewness ($s_\mu$) across benchmarks. Solid lines mark the posterior medians; dashed lines mark $\alpha_\mu = 2$, at which the Harvey curve reduces to a logistic, and $s_\mu = 0$.

Mazur Writing (0.860). Eighteen benchmarks have their asymptote at 1, thirteen pinned by a known ceiling and five by a recorded perfect score. For benchmarks with few observations, the interval is roughly the spread of the population distribution, since the pooling of Appendix A.5 then supplies most of the information. MultiChallenge and WMDP Chemistry have the widest intervals, as their best scores are still well below the population mean of the asymptotes.

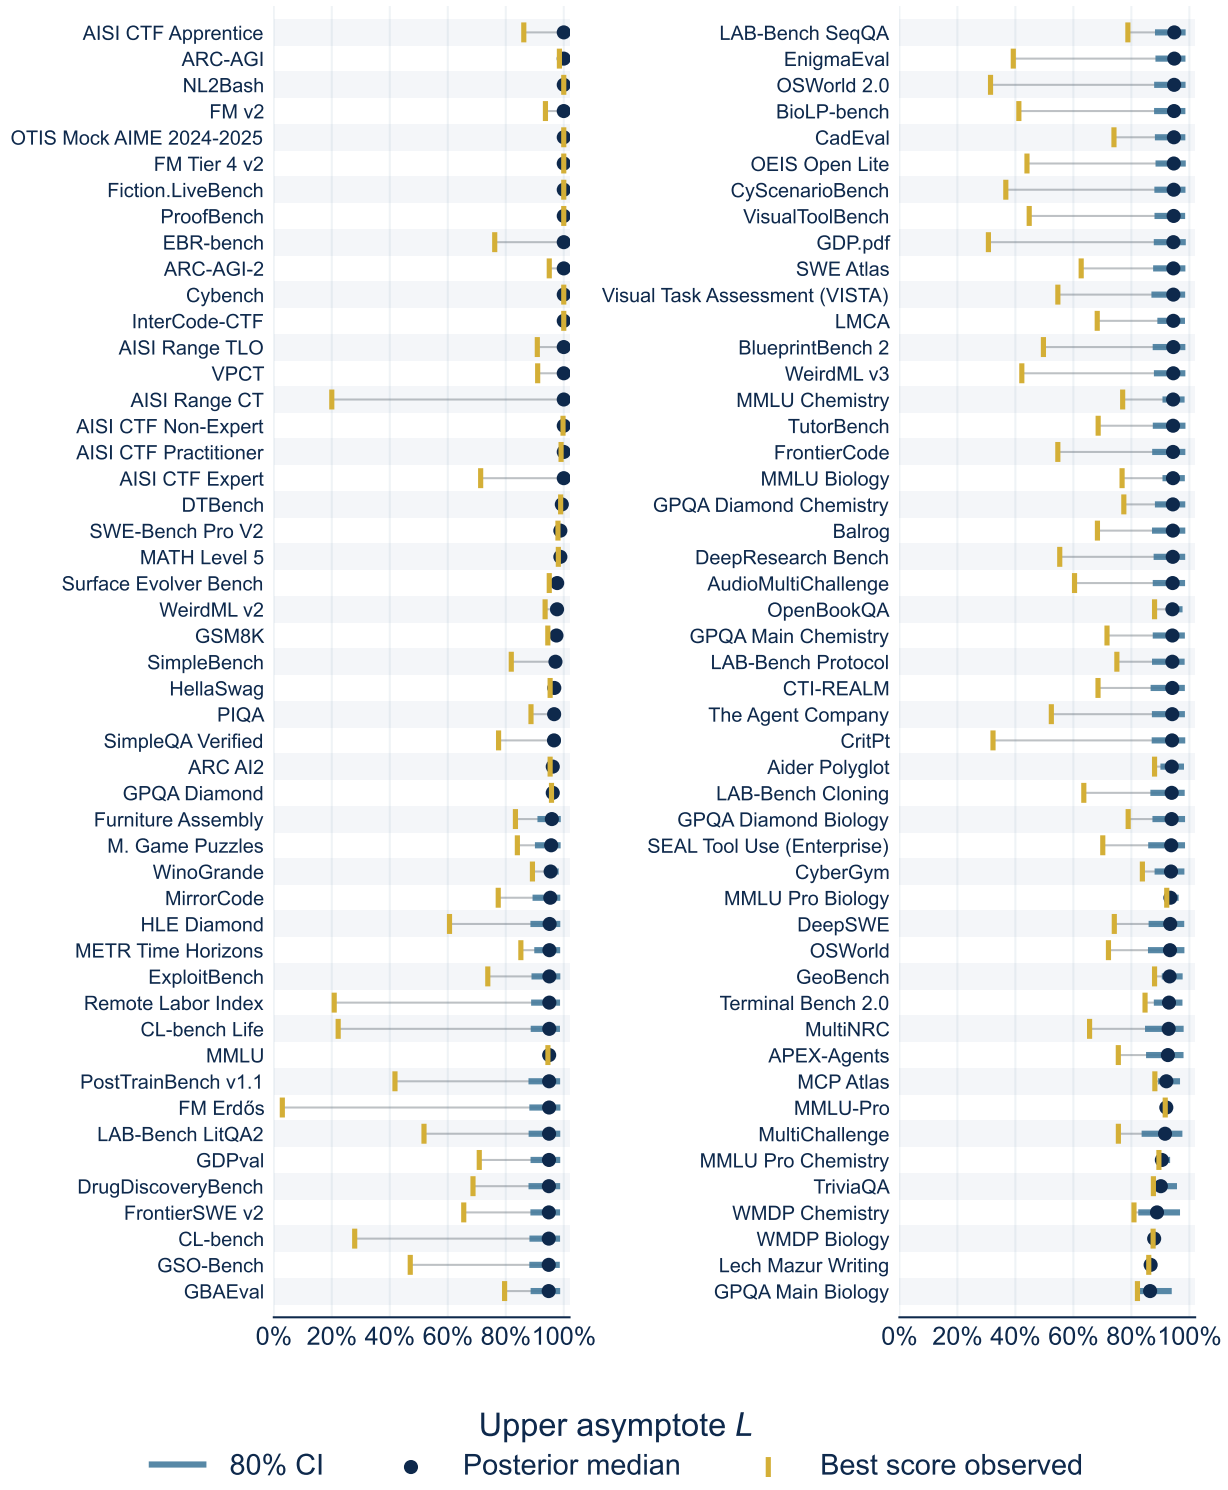


Figure 8: **Upper asymptote estimates by benchmark.** Posterior median and 80% credible interval of the upper asymptote L_i for each benchmark, ordered by median from the top of the left panel to the bottom of the right one, with the best score observed up to October 2026 marked by a vertical tick and joined to the median by a thin line. The hierarchical structure allows partial pooling, so benchmarks with little data are shrunk toward the population distribution.

E Benchmark Details

E.1 Included Benchmarks

Supp. Table 8 lists all 98 benchmarks with their category, observation count, lower bound, best score to date, posterior asymptote and projected saturation date. What each measures:

Core AGI progress and general reasoning. ARC-AGI and ARC-AGI-2 test how well a system picks up new abstract concepts from a few examples (Chollet, 2019; ARC Prize, 2019, 2025). The Remote Labor Index (Mazeika et al., 2025) counts the real freelance-platform contracts that a model completes. APEX-Agents (Vidgen et al., 2026) grades banking, consulting and law tasks of about two hours in realistic file environments, and in GDPval (Patwardhan et al., 2025) experts compare AI deliverables with professional work on tasks of about seven hours, across 44 U.S. occupations. Among general reasoning benchmarks, SimpleBench (Philip and Hemang, 2024) targets reasoning that humans find easy and Balrog (Paglieri et al., 2025) planning in game environments. EnigmaEval (Wang et al., 2025a) is made of long multimodal puzzles. Mystery Game Puzzles (Epoch AI, 2026c) asks for the best move in positions from an undisclosed puzzle variant of a well-known game. EBR-bench (Epoch AI, 2026a) measures how much a model improves at an unfamiliar campaign game (Earthborne Rangers) over repeated playthroughs with note-taking. CL-bench and CL-bench Life (Dou et al., 2026) test the learning of new knowledge, rules and procedures from task-specific contexts, professional in the first and from everyday life in the second. DTBench (Oesterheld et al., 2024) covers decision-theoretic reasoning in Newcomb-like problems. LMCA (Cooper et al., 2026) asks models to rate informal arguments in conceptual fields such as philosophy and AI safety, and scores them by correlation with expert ratings.

Domain-specific questions and mathematics. ARC AI2 (Clark et al., 2018) contains student-level science questions and MMLU (Hendrycks et al., 2020) four-option questions in 57 subjects, with MMLU-Pro (Wang et al., 2024) as a harder ten-option version. GPQA Diamond (Rein et al., 2023) has graduate-level science questions written to resist web search, and SimpleQA Verified (Wei et al., 2024; Haas et al., 2025) tests factual recall with verified answers. HLE Diamond is a cleaned and frozen subset of 1,000 expert-written questions from Humanity's Last Exam (Phan et al., 2025), across disciplines. CritPt (Zhu et al., 2025) poses unpublished research-level physics challenges, and GDP.pdf (Garre et al., 2026) questions grounded in real PDF documents from professional workflows. In mathematics, GSM8K (Cobbe et al., 2021) covers grade-school arithmetic, while MATH Level 5 (Hendrycks et al., 2021) and OTIS Mock AIME 2024–2025 (Epoch AI, 2025e) contain competition problems. FrontierMath Tiers 1–3 v2 and Tier 4 v2 (Epoch AI, 2025d) contain handcrafted problems designed to take expert mathematicians hours, Tier 4 being the hardest. FrontierMath Erdős (Adamczewski and Bloom, 2026) and OEIS Open Lite (Adamczewski, 2026) ask for open problems (Erdős problems and OEIS conjectures) to be resolved with a machine-checked Lean proof. ProofBench (Ravi et al., 2026) requires Lean 4 proofs of advanced problems that compile, with no partial credit.

Autonomous software engineering and agentic computer use. Aider Polyglot (Aider, 2024) is made of Exercism coding exercises in six languages. SWE-Bench Pro V2 (Jimenez et al., 2024; Deng et al., 2025) and GSO-Bench (Shetty et al., 2025) use real repositories. DeepSWE (Datacurve, 2026) has original tasks in 91 repositories and five languages, checked by hand-written behavioral verifiers, and FrontierCode (Cognition, 2026) asks whether open-source maintainers would accept the generated changes. FrontierSWE v2 (Proximal Labs, 2026) has 34 tasks at the limit of what humans can do, each with a 20-hour budget. SWE Atlas (Scale AI, 2026) asks questions about codebases and has models write tests and refactor code. MirrorCode (Adamczewski et al., 2026) shows how a program behaves, not its source, and asks for a reimplementation. GBAEval (Mechanize, 2026) asks for a Game Boy Advance emulator written from scratch in 24 hours and graded against a reference emulator, and Surface Evolver Bench (Henon, 2026) for physical simulations of liquid surfaces in a specialized solver. We score METR Time Horizons (Kwa et al., 2025) as the weighted mean success rate on METR's SWAA, HCAST and RE-Bench tasks, which measure autonomy on software and research tasks of increasing length. WeirdML v2 and v3 (Ihle, 2025) and PostTrainBench v1.1 (Rank et al., 2026) test machine-learning research and post-training, v3 being a new set of agentic tasks. For agentic computer use, Terminal Bench 2.0 (Merrill et al., 2026) and OSWorld (Xie et al., 2024) test the autonomous operation of a shell and of a desktop, and OSWorld 2.0 (Yuan et al., 2026) long professional desktop workflows that take skilled users hours. The Agent Company (Xu et al., 2025) reproduces the day-to-day work of a software company. MCP Atlas (Bandi et al., 2026) evaluates tool use via the Model Context Protocol, SEAL Tool Use (Enterprise) (Nath et al., 2025) multi-step reasoning with enterprise tools. DeepResearch Bench (FutureSearch) (FutureSearch et al., 2025) poses research questions to be answered from a frozen snapshot of the web.

Cyber. Cybench (Zhang et al., 2025) and InterCode-CTF (Yang et al., 2023) test vulnerability discovery and exploitation in capture-the-flag challenges, InterCode-CTF with picoCTF tasks in an interactive environment. The UK AI Security Institute grades the same skill with CTF suites at four difficulty levels (Technical Non-Expert, Apprentice, Practitioner, Expert), and uses two cyber ranges, The Last Ones and Cooling Tower, for multi-step attacks on a simulated network and on operational technology (UK AI Security Institute, 2026). CyScenarioBench (Irregular, 2026) asks the agent to plan and execute multi-stage attack scenarios. In CyberGym (Wang et al., 2025b), it must reproduce real-world vulnerabilities with proof-of-concept exploits. ExploitBench (Lee and Brumley, 2026) scores how far it gets in the V8 JavaScript engine, from reaching vulnerable code up to arbitrary code execution. CTI-REALM (Chakraborty et al., 2026) turns threat-intelligence reports into detection rules, and NL2Bash (Lin et al., 2018) turns natural-language requests into Bash commands.

Biology and chemistry. WMDP Biology and Chemistry (Li et al., 2024a) use proxy questions for knowledge relevant to weapons of mass destruction. The LAB-Bench subsets (Cloning, LitQA2, Protocol, SeqQA) (Laurent et al., 2024) and BioLP-bench (Dev et al., 2025) cover laboratory skills: understanding protocols, searching the literature, analyzing sequences. In DrugDiscoveryBench (Akyürek et al., 2026), coding agents work on early-stage drug discovery, from target identification to structure–activity analysis, and every task has a single verifiable answer. The biology and chemistry subsets of MMLU (Hendrycks et al., 2020), MMLU Pro (Wang et al., 2024), GPQA Diamond and GPQA Main (Rein et al., 2023) isolate these two disciplines.

Remaining categories. Trivia and commonsense QA covers OpenBookQA (Mihaylov et al., 2018) (elementary science with an open book of facts), HellaSwag (Zellers et al., 2019) (sentence continuation), PIQA (Bisk et al., 2020) (physical commonsense), WinoGrande (Sakaguchi et al., 2020) (Winograd-style pronoun resolution) and TriviaQA (Joshi et al., 2017) (open-ended trivia questions); language and writing, Lech Mazur Writing (Mazur, 2026) (judged short creative fiction), Fiction.LiveBench (Epoch AI, 2025c) (long-context comprehension of stories), TutorBench (Srinivasa et al., 2025) (tutoring), MultiChallenge (Sirdeshmukh et al., 2025) (realistic multi-turn conversations) and MultiNRC (Fabbri et al., 2025) (native multilingual reasoning); multimodal understanding, GeoBench (ccmdi, 2026) (geolocation of images), CadEval (Epoch AI, 2025a) (parametric 3D designs), Visual Task Assessment (VISTA) (Scale AI, 2025a) (rubric-graded visual tasks), VPCT (Brower, 2025) (visual physics comprehension), AudioMultiChallenge (Gosai et al., 2025) (multi-turn spoken dialogue), VisualToolBench (Guo et al., 2025) (tool-enabled image perception and transformation), BlueprintBench 2 (Petersson et al., 2025) (floor plans reconstructed from apartment photographs) and Furniture Assembly (Epoch AI, 2026b) (finding the mistaken step in photographs of furniture builds against the official manual).

Table 8: **Per-benchmark summary.** Lower bound ℓ is random-chance performance where it is well defined (see Appendix E.2); L is the posterior upper asymptote (median and 80% CI); the saturation date is the posterior median of the date at which the trajectory reaches 95% of the score range, with its 80% credible interval. n counts frontier observations used for fitting, and "Best obs." is the highest score recorded up to October 2026. Rows are grouped by capability category.

Benchmark	n	ℓ	Best obs.	L [80% CI]	Saturation date [80% CI]
Advanced Language and Writing					
Fiction.LiveBench	15	0.000	1.00	1.00 [1.00, 1.00]	2025-06 [2025-05, 2025-08]
Lech Mazur Writing	16	0.000	0.86	0.87 [0.86, 0.88]	2024-11 [2024-10, 2025-02]
MultiChallenge	17	0.000	0.76	0.92 [0.84, 0.98]	2027-03 [2026-07, 2027-11]
MultiNRC	15	0.000	0.66	0.93 [0.85, 0.98]	2027-04 [2026-10, 2027-12]
TutorBench	10	0.000	0.69	0.94 [0.87, 0.99]	2028-02 [2027-05, 2029-04]
Agentic Computer Use					
DeepResearch Bench	12	0.000	0.55	0.94 [0.88, 0.99]	2029-08 [2028-06, 2031-05]
MCP Atlas	20	0.000	0.88	0.92 [0.89, 0.97]	2026-06 [2026-04, 2026-09]
OSWorld	9	0.000	0.72	0.93 [0.86, 0.98]	2026-09 [2026-05, 2027-01]
OSWorld 2.0	7	0.000	0.31	0.95 [0.88, 0.99]	2027-09 [2027-04, 2028-12]
SEAL Tool Use (Enterprise)	12	0.000	0.70	0.94 [0.86, 0.99]	2027-11 [2026-07, 2029-08]
Terminal Bench 2.0	22	0.000	0.85	0.93 [0.88, 0.98]	2026-05 [2026-03, 2026-07]
The Agent Company	12	0.000	0.52	0.94 [0.87, 0.98]	2026-11 [2026-06, 2027-08]
Autonomous SWE					
Aider Polyglot	20	0.000	0.88	0.94 [0.90, 0.98]	2025-07 [2025-05, 2025-09]

Table 8, continued

Benchmark	n	ℓ	Best obs.	L [80% CI]	Saturation date [80% CI]
DeepSWE	12	0.000	0.74	0.93 [0.86, 0.98]	2026-11 [2026-09, 2027-02]
FrontierCode	13	0.000	0.55	0.94 [0.87, 0.99]	2027-10 [2027-05, 2028-05]
FrontierSWE v2	9	0.000	0.66	0.95 [0.88, 0.99]	2027-06 [2027-01, 2028-10]
GBAEval	10	0.000	0.80	0.95 [0.89, 0.99]	2026-10 [2026-08, 2027-01]
GSO-Bench	18	0.000	0.47	0.95 [0.88, 0.99]	2027-04 [2027-01, 2027-09]
METR Time Horizons	31	0.050	0.85	0.95 [0.90, 0.99]	2026-09 [2026-06, 2027-01]
MirrorCode	9	0.000	0.77	0.95 [0.89, 0.99]	2027-02 [2026-11, 2027-08]
PostTrainBench v1.1	9	0.075	0.42	0.95 [0.88, 0.99]	2027-12 [2027-05, 2029-07]
SWE Atlas	15	0.000	0.63	0.94 [0.88, 0.99]	2027-06 [2027-03, 2027-10]
SWE-Bench Pro V2	8	0.000	0.98	0.99 [0.98, 1.00]	2026-08 [2026-06, 2026-09]
Surface Evolver Bench	11	0.000	0.95	0.98 [0.96, 0.99]	2026-11 [2026-07, 2027-06]
WeirdML v2	33	0.000	0.94	0.98 [0.95, 0.99]	2026-07 [2026-05, 2026-10]
WeirdML v3	8	0.000	0.42	0.94 [0.88, 0.99]	2027-12 [2027-06, 2029-05]
Biology					
BioLP-bench	17	0.000	0.41	0.95 [0.88, 0.99]	2026-08 [2026-02, 2027-04]
DrugDiscoveryBench	13	0.000	0.69	0.95 [0.88, 0.99]	2027-06 [2027-02, 2027-10]
GPQA Diamond Biology	11	0.250	0.79	0.94 [0.87, 0.99]	2025-10 [2025-03, 2026-09]
GPQA Main Biology	19	0.250	0.82	0.86 [0.83, 0.94]	2025-06 [2025-02, 2026-02]
LAB-Bench Cloning	14	0.200	0.64	0.94 [0.87, 0.98]	2026-06 [2025-12, 2027-04]
LAB-Bench LitQA2	12	0.200	0.52	0.95 [0.88, 0.99]	2027-12 [2026-11, 2029-12]
LAB-Bench Protocol	13	0.200	0.75	0.94 [0.87, 0.98]	2025-11 [2025-04, 2027-01]
LAB-Bench SeqQA	16	0.200	0.79	0.95 [0.88, 0.99]	2026-03 [2025-10, 2026-10]
MMLU Biology	14	0.250	0.77	0.94 [0.91, 0.98]	2026-06 [2025-10, 2027-08]
MMLU Pro Biology	18	0.100	0.92	0.93 [0.92, 0.96]	2024-04 [2024-01, 2024-09]
WMDP Biology	18	0.250	0.88	0.88 [0.88, 0.89]	2024-01 [2023-11, 2024-03]
Chemistry					
GPQA Diamond Chemistry	20	0.250	0.77	0.94 [0.88, 0.99]	2025-11 [2025-07, 2026-04]
GPQA Main Chemistry	21	0.250	0.72	0.94 [0.87, 0.98]	2026-01 [2025-10, 2026-08]
MMLU Chemistry	18	0.250	0.77	0.94 [0.91, 0.98]	2026-01 [2025-08, 2026-09]
MMLU Pro Chemistry	21	0.100	0.90	0.91 [0.90, 0.93]	2024-08 [2024-06, 2024-10]
WMDP Chemistry	20	0.250	0.81	0.89 [0.82, 0.97]	2025-06 [2024-10, 2026-04]
Core AGI Progress					
APEX-Agents	18	0.000	0.76	0.93 [0.85, 0.98]	2027-03 [2026-12, 2027-06]
ARC-AGI	29	0.000	0.98	1.00 [1.00, 1.00]	2026-01 [2025-12, 2026-02]
ARC-AGI-2	31	0.000	0.95	1.00 [1.00, 1.00]	2026-06 [2026-06, 2026-07]
GDPval	10	0.000	0.71	0.95 [0.89, 0.99]	2026-08 [2026-05, 2026-12]
Remote Labor Index	15	0.000	0.21	0.95 [0.89, 0.99]	2028-01 [2027-09, 2028-09]
Cyber					
AISI CTF Suites - Apprentice	17	0.000	0.86	1.00 [1.00, 1.00]	2026-09 [2026-07, 2026-12]
AISI CTF Suites - Expert	8	0.000	0.71	1.00 [1.00, 1.00]	2027-01 [2026-10, 2027-07]
AISI CTF Suites - Practitioner	8	0.000	0.99	1.00 [1.00, 1.00]	2026-02 [2026-01, 2026-03]
AISI CTF Suites - Technical Non-Expert	13	0.000	1.00	1.00 [1.00, 1.00]	2025-05 [2025-03, 2025-08]
AISI Cyber Ranges - Cooling Tower	6	0.000	0.20	1.00 [1.00, 1.00]	2027-02 [2026-10, 2027-11]
AISI Cyber Ranges - The Last Ones	13	0.000	0.91	1.00 [1.00, 1.00]	2026-08 [2026-07, 2026-10]
CTI-REALM	10	0.000	0.69	0.94 [0.87, 0.98]	2027-06 [2026-11, 2028-03]
CyScenarioBench	9	0.000	0.37	0.95 [0.88, 0.99]	2027-05 [2027-01, 2028-04]
Cybench	17	0.000	1.00	1.00 [1.00, 1.00]	2026-02 [2026-01, 2026-04]
CyberGym	12	0.000	0.84	0.94 [0.88, 0.98]	2026-06 [2026-04, 2026-08]
ExploitBench	8	0.000	0.74	0.95 [0.89, 0.99]	2027-01 [2026-09, 2028-02]
InterCode-CTF	12	0.000	1.00	1.00 [1.00, 1.00]	2025-01 [2024-11, 2025-03]
NL2Bash	13	0.000	1.00	1.00 [1.00, 1.00]	2025-06 [2025-01, 2025-12]
Domain Specific Questions					
ARC AI2	17	0.250	0.95	0.96 [0.95, 0.98]	2023-06 [2023-02, 2023-10]
CritPt	25	0.000	0.32	0.94 [0.87, 0.99]	2028-06 [2027-11, 2029-02]
GDP.pdf	11	0.000	0.31	0.94 [0.88, 0.99]	2028-05 [2027-10, 2029-05]
GPQA Diamond	45	0.250	0.96	0.96 [0.96, 0.97]	2025-10 [2025-09, 2025-12]
HLE Diamond	8	0.048	0.61	0.95 [0.89, 0.99]	2027-04 [2027-01, 2028-01]
MMLU	37	0.250	0.95	0.95 [0.95, 0.96]	2024-03 [2023-11, 2024-06]
MMLU-Pro	28	0.100	0.92	0.92 [0.92, 0.93]	2025-05 [2025-04, 2025-07]

Table 8, continued

Benchmark	n	ℓ	Best obs.	L [80% CI]	Saturation date [80% CI]
SimpleQA Verified	20	0.000	0.78	0.97 [0.95, 0.99]	2027-06 [2027-02, 2027-11]
General Reasoning					
Balrog	17	0.000	0.68	0.94 [0.87, 0.99]	2028-03 [2027-07, 2029-01]
CL-bench	11	0.000	0.28	0.95 [0.88, 0.99]	2030-01 [2028-07, 2032-12]
CL-bench Life	9	0.000	0.22	0.95 [0.89, 0.99]	2028-09 [2027-09, 2030-12]
DTBench	39	0.400	0.99	0.99 [0.99, 1.00]	2026-01 [2025-12, 2026-03]
EBR-bench	16	0.000	0.76	1.00 [1.00, 1.00]	2027-03 [2027-01, 2027-06]
EnigmaEval	21	0.000	0.39	0.95 [0.88, 0.99]	2027-12 [2027-06, 2028-12]
LMCA	34	0.000	0.68	0.94 [0.89, 0.98]	2027-10 [2027-06, 2028-03]
Mystery Game Puzzles	23	0.092	0.84	0.96 [0.90, 0.99]	2027-03 [2027-01, 2027-06]
SimpleBench	23	0.167	0.82	0.97 [0.96, 0.99]	2027-04 [2026-12, 2027-08]
Mathematics					
FrontierMath Erdős	4	0.000	0.03	0.95 [0.88, 0.99]	2029-07 [2027-11, 2032-10]
FrontierMath Tier 4 v2	21	0.000	1.00	1.00 [1.00, 1.00]	2026-08 [2026-07, 2026-09]
FrontierMath Tiers 1-3 v2	26	0.000	0.94	1.00 [1.00, 1.00]	2026-08 [2026-07, 2026-10]
GSM8K	16	0.000	0.94	0.97 [0.97, 0.99]	2023-09 [2023-06, 2023-12]
MATH Level 5	30	0.000	0.98	0.99 [0.98, 1.00]	2024-12 [2024-11, 2025-01]
OEIS Open Lite	5	0.000	0.44	0.95 [0.88, 0.99]	2028-05 [2027-05, 2032-03]
OTIS Mock AIME 2024-2025	45	0.001	1.00	1.00 [1.00, 1.00]	2025-09 [2025-08, 2025-10]
ProofBench	20	0.000	1.00	1.00 [1.00, 1.00]	2026-08 [2026-07, 2026-09]
Multimodal Understanding					
AudioMultiChallenge	12	0.000	0.60	0.94 [0.87, 0.99]	2027-11 [2027-05, 2028-11]
BlueprintBench 2	12	0.000	0.50	0.94 [0.87, 0.99]	2027-09 [2027-05, 2028-05]
CadEval	11	0.000	0.74	0.95 [0.88, 0.99]	2025-11 [2025-08, 2026-03]
Furniture Assembly	15	0.300	0.83	0.96 [0.91, 0.99]	2026-12 [2026-11, 2027-02]
GeoBench	12	0.000	0.88	0.93 [0.90, 0.98]	2025-07 [2025-03, 2026-01]
VPCT	15	0.333	0.91	1.00 [1.00, 1.00]	2026-06 [2026-03, 2026-10]
Visual Task Assessment (VISTA)	19	0.000	0.55	0.94 [0.87, 0.99]	2028-08 [2027-09, 2029-11]
VisualToolBench	12	0.000	0.45	0.95 [0.88, 0.99]	2027-11 [2027-05, 2028-08]
Trivia & Commonsense QA					
HellaSwag	19	0.250	0.95	0.97 [0.96, 0.99]	2023-08 [2023-01, 2024-04]
OpenBookQA	14	0.250	0.88	0.94 [0.92, 0.98]	2024-03 [2023-10, 2024-10]
PIQA	18	0.500	0.89	0.97 [0.95, 0.99]	2027-08 [2026-05, 2029-05]
TriviaQA	11	0.127	0.88	0.90 [0.88, 0.96]	2022-10 [2022-06, 2023-08]
WinoGrande	13	0.500	0.89	0.95 [0.94, 0.98]	2024-05 [2023-09, 2025-05]

E.2 Lower Bounds

Supp. Table 8 gives the lower bound of each benchmark, and 32 of the 98 are above zero. For 25 of them the bound follows from the answer format, one over the number of options for multiple choice: 0.50 for WinoGrande and PIQA, 0.333 for VPCT, 0.25 for GPQA Diamond, MMLU, their biology and chemistry subsets and those of GPQA Main, WMDP Biology and Chemistry, ARC AI2, HellaSwag and OpenBookQA, 0.20 for the four LAB-Bench subsets (five options on average), 0.167 for SimpleBench and 0.10 for MMLU-Pro and its two subsets. OTIS Mock AIME 2024–2025 takes 0.001, its answers being integers from 000 to 999. Four bounds are Epoch AI's published chance baselines (DTBench 0.40, Furniture Assembly 0.30, Mystery Game Puzzles 0.0922, HLE Diamond 0.048, the last inherited from Humanity's Last Exam). The remaining three need a word of explanation. PostTrainBench reports the average benchmark score of the post-trained models, and an agent that does nothing leaves the base models at 7.5%, so we use 0.075. TriviaQA takes 0.127, the random-entity baseline of its authors. For METR Time Horizons we take 0.05, the share of multiple-choice sub-tasks in the suite.

A zero bound is justified for 65 of the other 66 benchmarks. Free-response, agentic and exact-match tasks give nothing to a random output; DeepResearch Bench, for example, grades answers on precision, recall, F1 or exact correctness. On GeoBench, a random guess finds the right country about once in two hundred tries. Some scores are built to be zero at chance: GDPval counts wins plus ties against human experts, LMCA correlates with expert ratings, and BlueprintBench 2 is normalized so that random output scores zero. Lech Mazur Writing is the exception, as an LLM judge grades short fiction and no one has measured a random

baseline. We still set its bound to zero, since we expect chance-level output to score below 10%, and the approximation is small.

E.3 Cross-Benchmark Conditional Dependence

The model treats trajectories as conditionally independent given the shared hyperpriors, and several benchmarks share a source dataset or a creator. We pair residuals $y_i(t) - \mathbb{E}[\mu_i(t)]$ by model release and correlate them across benchmarks over the releases they have in common. The correlations are positive throughout: +0.30 on average between unrelated benchmarks, and +0.47 within the twelve lineages built on a common dataset or maintainer (the GPQA Diamond, GPQA Main, MMLU, MMLU Pro, LAB-Bench, ARC-AGI, FrontierMath, WMDP, MultiChallenge and CL-bench families, and the UK AI Security Institute's capture-the-flag suites and cyber ranges).

The baseline level is mostly a general-capability effect, where strong models sit above trend on everything they are evaluated on. Model identity indeed accounts for 27 to 36% of residual variance (adjusted and raw R^2), so this component adds shared noise around the fitted curves and impacts the intervals.

Our intervals thus treat our 98 benchmarks as carrying a bit more independent information than they do, but reducing the dataset to only one representative per lineage (leaving 78 benchmarks) gives a similar projected saturation proportion, of 97.4% (80% CI: $[94.9, 98.7]$) against 98.0% ($[95.9, 99.0]$) on all 98.

F Human Performance Baselines

Supp. Table 9 summarizes human performance baselines used in our analysis. We categorize human performance into four main groups based on expertise level, with committee and high-school variants:

- **Average Human**: Crowdworkers (e.g., MTurk) or non-specialized participants
- **Skilled Generalist**: Individuals with advanced education but not in the target domain (e.g., PhD students in unrelated fields, skilled professionals)
- **Domain Expert**: PhD-level specialists in the relevant domain or expert professionals
- **Top Performer**: Elite performers (e.g., top 5% test takers or best result)

Committee scores represent an aggregation of majority votes or average team scores across human participants. The *High School Qualifier* and *High School Top Performer* categories specifically represent a cohort of students in a specialized training program.

Coverage and quality. Coverage is uneven, good for the dual-use biology and chemistry benchmarks, where RAND elicited expert performance on the same items (Dev et al., 2025), limited to Cybench and InterCode-CTF among the cybersecurity benchmarks, and absent for writing benchmarks and for most agentic and software engineering ones. The baselines also differ in kind (time limits, incentives, tool access, etc.) and sample sizes are rarely large. See Wei et al. (2025) for a call for more rigorous and better-documented baselines to support human vs. AI comparison, as well as a practical checklist to structure human metadata.

Table 9: **Human performance baselines by benchmark.** Scores represent accuracy unless otherwise noted. Committee scores use majority votes or average team scores across human participants. Details regarding each human baseline (population, protocol, source) are available with the data in the code repository.

Benchmark	Human Group	Score	Source
Domain Specific Questions			
GPQA Diamond	Domain Expert	69.7%	OpenAI (2024)
GPQA Diamond	Skilled Generalist	21.9%	Rein et al. (2023)
MMLU	High School Top Performer	89.8%	Hendrycks et al. (2020)
MMLU	Average Human	34.5%	Hendrycks et al. (2020)
SimpleQA Verified	Skilled Generalist	94.4%	Wei et al. (2024)
General Reasoning			
DTBench	Domain Expert	98.7%	Oesterheld et al. (2024)
EBR-bench	Top Performer	100%	Epoch AI (2024)
EBR-bench	Average Human	52.8%	Epoch AI (2024)

Table 9, continued

Benchmark	Human Group	Score	Source
LMCA	Domain Expert	87%	Cooper et al. (2026)
SimpleBench	Top Performer	95.4%	Philip and Hemang (2024)
SimpleBench	Average Human	83.7%	Philip and Hemang (2024)
Mathematics			
FrontierMath Erdős	Domain Expert	0%	Adamczewski and Bloom (2026)
FrontierMath Tiers 1–3 v2	Committee of Domain Experts	35%	Epoch AI (2025d)
GSM8K	Skilled Generalist	96.8%	Li et al. (2024b)
MATH Level 5	Domain Expert	90%	Hendrycks et al. (2021)
MATH Level 5	Skilled Generalist	40%	Hendrycks et al. (2021)
OTIS Mock AIME	High School Top Performer	93%	Chen (2025)
OTIS Mock AIME	High School Qualifier	53%	Chen (2025)
Core AGI Progress			
ARC-AGI	Committee of Average Humans	98%	ARC Prize (2019)
ARC-AGI	Skilled Generalist	98%	ARC Prize (2019)
ARC-AGI	Average Human	77%	ARC Prize (2019)
ARC-AGI-2	Committee of Average Humans	100%	ARC Prize (2025)
ARC-AGI-2	Average Human	60%	ARC Prize (2025)
Agentic Computer Use			
OSWorld	Skilled Generalist	72.4%	Xie et al. (2024)
The Agent Company	Domain Expert	75%	Wang et al. (2025c)
Autonomous SWE			
METR Time Horizons	Committee of Domain Experts	83%	Kwa et al. (2025)
Cyber			
Cybench	Committee of Domain Experts	75%	CTF scoreboards
InterCode-CTF	Domain Expert	95.6%	Payne et al. (2026)
Biology (Dual-Use)			
BioLP-bench	Domain Expert	38.4%	Dev et al. (2025)
GPQA Diamond Biology	Domain Expert	83.1%	Dev et al. (2025)
GPQA Diamond Biology	Skilled Generalist	21.9%	Rein et al. (2023)
GPQA Main Biology	Domain Expert	66.7%	Dev et al. (2025)
GPQA Main Biology	Skilled Generalist	43.2%	Rein et al. (2023)
LAB-Bench Cloning	Domain Expert	60%	Dev et al. (2025)
LAB-Bench LitQA2	Domain Expert	70%	Dev et al. (2025)
LAB-Bench Protocol	Domain Expert	79%	Dev et al. (2025)
LAB-Bench SeqQA	Domain Expert	78%	Dev et al. (2025)
MMLU Biology	High School Top Performer	89.8%	Dev et al. (2025)
WMDP Biology	Domain Expert	60.5%	Dev et al. (2025)
Chemistry (Dual-Use)			
GPQA Diamond Chemistry	Domain Expert	83.1%	Dev et al. (2025)
GPQA Diamond Chemistry	Skilled Generalist	21.9%	Rein et al. (2023)
GPQA Main Chemistry	Domain Expert	72%	Dev et al. (2025)
GPQA Main Chemistry	Skilled Generalist	31.4%	Rein et al. (2023)
MMLU Chemistry	High School Top Performer	89.8%	Dev et al. (2025)
WMDP Chemistry	Domain Expert	43.3%	Dev et al. (2025)
Trivia & Commonsense QA			
HellaSwag	Committee of Average Humans	95.6%	Zellers et al. (2019)
OpenBookQA	Average Human	92%	Mihaylov et al. (2018)
PIQA	Committee of Skilled Generalists	94.9%	Bisk et al. (2020)
TriviaQA	Average Human	79.7%	Joshi et al. (2017)
WinoGrande	Committee of Average Humans	94%	Sakaguchi et al. (2020)
Multimodal Understanding			
BlueprintBench 2	Skilled Generalist	58.6%	Epoch AI (2024)
Furniture Assembly	Committee of Average Humans	84%	Epoch AI (2026b)
Furniture Assembly	Average Human	63%	Epoch AI (2026b)
GeoBench	Top Performer	90%	ccmdi (2026)
VISTA	Skilled Generalist	55.4%	Scale AI (2025a)

VPCT	Average Human	100%	Brower (2025)

Interpretation notes. Human baselines change a lot with expertise and test conditions. On GPQA Diamond, skilled generalists score 22% and domain experts 70%, a gap of nearly 50 points that shows how much the difficulty of the benchmark depends on domain knowledge. On MATH Level 5, a CS PhD student scores 40% and an elite mathematician 90%, as competition mathematics is a skill of its own.

When frontier models exceed domain-expert performance, as they now do on GPQA Diamond, WMDP Biology and Chemistry, Cybench and InterCode-CTF among others, they cross a meaningful capability threshold. A benchmark score may still not translate into real-world task completion, and human baselines have their own limits, such as small samples, varied incentives and ceiling effects.